

Diagnosing and Correcting NaN Cascades in Learned Dirichlet–Neumann Operator Rollouts

Internal technical report

9 July 2026

Contents

1	Reader’s guide and executive conclusions	3
2	Problem and numerical setting	4
2.1	Reference and surrogate operators	5
2.2	The integrating-factor GL2 scheme	5
2.3	What counts as success	5
3	What the experimental sequence established	6
3.1	Baseline numbers and checkpoint conventions	6
3.2	Phase I: establish a trustworthy evaluation	7
3.3	Phase II: test the data-coverage explanation	7
3.4	Phase III: test the integrator explanation	7
3.5	Phase IV: suppress the observed spectrum	7
4	Frame-by-frame failure anatomy	9
5	Hamiltonian inconsistency is the root defect	10
5.1	The measured defect	10
6	The missing spectral null form	11
7	Aliasing and the role of the integrator	12
8	C16: the resulting model and objective	13
8.1	Residual parameterization and its invariants	13
8.2	Finite-secant Hadamard loss	14
8.3	Optimizer calibration	14
8.4	What C16 guarantees, and what it does not	15
9	Current experiment and acceptance criteria	15
9.1	Planned post-training decision sequence	16

10 Alternative explanations, objections, and falsifiers	17
10.1 Could the data still contain subtle corruption?	17
10.2 Does small value error plus large derivative error necessarily cause a NaN?	17
10.3 Could GL2 still be the primary issue?	17
10.4 Why not learn the whole Hamiltonian and differentiate it?	17
10.5 Why not impose a global Lipschitz or spectral norm bound?	18
10.6 What result would falsify the C16 mechanism?	18
A Chronological experiment map	18
B Implementation and correctness checks	20

Abstract

Accurate one-step approximation of a water-wave Dirichlet–Neumann operator (DNO) did not prevent rare catastrophic failures in long symplectic rollouts. Frame-aligned trajectory analysis shows why: the learned operator is accurate on exact solution states but has an incorrect derivative transverse to the solution manifold. It consequently violates the Hadamard shape-derivative identity, so the two equations advanced by the Gauss–Legendre integrator are not the canonical derivatives of one learned Hamiltonian. Small state errors then excite spurious same-sign spectral sidebands, quadratic aliasing amplifies the already established cascade, and the implicit stage iteration eventually loses contractivity. This report summarizes the experimental path that ruled out data corruption, time-step refinement, spectral caps, hard cutoffs, and targeted fine-tuning as primary cures. It motivates a clean retraining of $G_0 + G_1 + \mathcal{R}_\theta$, where the analytic first-order Craig–Sulem term is exact, $\mathcal{R}_\theta = O(\eta^2)$ is self-adjoint, and a randomized Hadamard loss directly constrains the missing shape derivative. Particular care is taken to separate what has been directly measured, what follows mathematically from the measured facts, and what remains a hypothesis to be tested by the new model.

1 Reader’s guide and executive conclusions

The project accumulated many experiments because the terminal symptom—a NaN inside an implicit time step—admits several superficially plausible explanations. A reader should not be expected to remember the experiment labels. The labels C1–C16 are therefore treated here only as reproducibility handles. The scientific argument is organized around five questions:

1. Is the reference data corrupt or missing the failing regime?
2. Does the learned DNO first make an error on an exact physical state, or does it react incorrectly to the small off-manifold state error created by its own rollout?
3. Is the cascade initiated by the DNO, by pseudospectral aliasing, or by nonconvergence of the implicit Gauss–Legendre stages?
4. Which exact property of a physical DNO is absent from the learned map?
5. What is the smallest architectural and loss-level correction that restores that property without suppressing legitimate high-frequency dynamics?

The present answers, with their evidentiary status, are:

Directly measured:

The failing models can have relative DNO value error of only 2×10^{-4} on the exact state while their response to the difference between the exact and surrogate states grows to a secant gain of 7–8. The induced response first grows in $k = 32, \dots, 64$, then contaminates $k = 64, \dots, 128$, and only afterward makes the quadratic right-hand side strongly aliased and the fixed-point stage solve noncontractive.

Directly measured structural defect:

At the same exact states, the learned operator violates the DNO Hadamard shape-derivative identity by relative $O(1)$, whereas the order-six reference implementation satisfies the identical discrete test to approximately 10^{-5} or better. This discrepancy is far larger than the DNO value error and is already present before visible blow-up.

Mathematical consequence:

If a learned DNO is self-adjoint in ξ but fails the shape identity, then the implemented formulas for η_t and ξ_t are not the two variational derivatives of one learned Hamiltonian. A symplectic integrator cannot restore a Hamiltonian structure absent from the vector field it is given.

Strong mechanistic evidence:

The first Craig–Sulem term has a same-sign Fourier null form that cancels the specific high-mode sidebands observed in the failures. The learned operator does not preserve this cancellation. Exact $G_0 + G_1$ gives 0/32 Tanaka NaNs but is not accurate enough at final time, strongly motivating it as a stable backbone rather than a complete surrogate.

Hypothesis under test:

A full-capacity model of the form $G_0 + G_1 + O(\eta^2)$, trained with a direct Hadamard-consistency loss, should retain the stable null form while learning the small higher-order correction. This is C16. Its structural invariants and short training smoke have passed; the decisive long-rollout evaluation has not yet been performed.

This hierarchy matters. We have compelling evidence for the diagnosis, but we do not yet have evidence that C16 reaches the operational target of zero NaNs. A negative C16 result would not erase the measured shape defect; it would mean that the particular parameterization, loss estimator, training budget, or some additional missing identity remains inadequate.

2 Problem and numerical setting

We consider a periodic one-dimensional horizontal domain of length 2π , with free surface $y = \eta(x, t)$, constant depth h , and an irrotational velocity potential ϕ in the fluid. Its trace at the free surface is $\xi(x, t) = \phi(x, \eta(x, t), t)$. The Dirichlet–Neumann operator is the map from this trace to the non-unit normal derivative,

$$G(\eta; h)\xi = (\phi_y - \eta_x \phi_x)|_{y=\eta}. \quad (1)$$

The definition incorporates the solution of Laplace’s equation in the entire fluid domain. It is therefore the expensive nonlocal component that the surrogate is intended to replace.

For surface elevation η and surface potential ξ , Zakharov’s canonical water-wave equations without surface tension may be written

$$\eta_t = G(\eta)\xi, \quad (2)$$

$$\xi_t = -g\eta - \frac{1}{2}\xi_x^2 + \frac{1}{2} \frac{(G(\eta)\xi + \eta_x \xi_x)^2}{1 + \eta_x^2}. \quad (3)$$

They are Hamilton’s equations for

$$H(\eta, \xi) = \frac{1}{2} \int_0^{2\pi} \xi G(\eta) \xi \, dx + \frac{g}{2} \int_0^{2\pi} \eta^2 \, dx, \quad \eta_t = \frac{\delta H}{\delta \xi}, \quad \xi_t = -\frac{\delta H}{\delta \eta}. \quad (4)$$

This equivalence uses both self-adjointness of $G(\eta)$ in its potential argument and its precise derivative with respect to the surface. Preserving only the former does not suffice.

The production surrogate replaces the expensive DNO G by a learned operator $\mathcal{G}_\theta$, while retaining the geometric formula for ξ_t . Time integration uses a two-stage Gauss–Legendre method with an integrating factor and four fixed-point iterations per step. The DNO model is linear in ξ and built from Craig–Sulem-shaped multiplier sandwiches.

2.1 Reference and surrogate operators

The reference DNO is a truncated Craig–Sulem series evaluated through order six with substantial Fourier zero padding. It is not an exact continuum oracle, but its convergence, self-adjointness, shape-identity residual, and long-rollout conservation are all sufficiently accurate to distinguish the effects discussed below. In particular, the reference shape defect is 10^{-5} or smaller at the failing states, versus $O(1)$ for the learned operator.

The principal baseline is an eight-block Craig–Sulem DNO network. Each block is linear in ξ : a depth-conditioned Fourier multiplier is applied to ξ , multiplied pointwise by learned features of η , and passed through another Fourier multiplier. This is a much stronger inductive bias than a generic FNO and produced excellent one-step validation error. Later checkpoints tied the input and output multipliers so that the learned residual is self-adjoint in ξ . Nothing in the original architecture, however, forces the correct derivative with respect to η .

2.2 The integrating-factor GL2 scheme

Following Xu and Guyenne’s spectral Hamiltonian method [1], split the semi-discrete vector field as

$$u_t = Au + N(u), \quad u = (\eta, \xi)^\top, \quad A = \begin{pmatrix} 0 & G_0 \\ -g & 0 \end{pmatrix}, \quad (5)$$

where $G_0 = |D| \tanh(h|D|)$. With $v(t) = e^{-tA}u(t)$, the stiff linear dispersion is treated exactly and the transformed equation is advanced with the two-stage, fourth-order Gauss–Legendre Runge–Kutta method. Its Butcher coefficients are

$$a_{11} = a_{22} = \frac{1}{4}, \quad a_{12} = \frac{1}{4} - \frac{\sqrt{3}}{6}, \quad a_{21} = \frac{1}{4} + \frac{\sqrt{3}}{6}, \quad b_1 = b_2 = \frac{1}{2}, \quad (6)$$

up to interchange of the two stage labels. The coupled stage equations are solved here with four fixed-point iterations.

Three logically distinct statements must not be conflated:

1. GL2 is a symplectic Runge–Kutta method for a canonical Hamiltonian vector field.
2. A finite number of Picard iterations only approximately solves the implicit stage equations; loss of contractivity can expose a bad state.
3. If the supplied surrogate vector field is not Hamiltonian, exact solution of the GL2 stages still cannot confer the missing Hamiltonian geometry on that field.

The investigation below identifies item 3 as the initiating issue and item 2 as a terminal failure mechanism.

2.3 What counts as success

The practical target is stricter than small one-step validation loss: across the 32-case Tanaka test and the 31 truth-valid Benjamin–Feir cases, the rollout must produce no nonfinite trajectories while retaining competitive median and tail error at final time.

We distinguish four metrics:

One-step DNO error

The relative error in $G(\eta)\xi$ when the model is evaluated on an exact reference state. This measures interpolation on the sampled truth manifold.

NaN rate

The fraction of truth-valid initial conditions for which any surrogate state becomes nonfinite. This is the primary hard stability endpoint.

Divergence rate

The fraction whose final relative state error exceeds the evaluation threshold, including but not limited to nonfinite cases. A finite but physically useless trajectory is not counted as a success.

Final-time error distribution

The median and 95th percentile relative L^2 error in η . Reporting only the median can hide a cap or filter that protects easy cases while producing a catastrophic tail.

The rollout harness stores the integrator state and spectral operations in double precision while leaving the neural network in its trained precision. This matters: an earlier single-precision harness imposed an artificial long-time error floor. Truth-invalid initial conditions are excluded using the reference solver’s own nonfiniteness or excessive Hamiltonian drift, so a surrogate is not penalized for a failed reference trajectory.

3 What the experimental sequence established

Table 2 condenses the main branches recorded in `EXPERIMENTS.md`. The sequence is informative because several apparently plausible interventions improved an intermediate diagnostic without eliminating the failure population.

3.1 Baseline numbers and checkpoint conventions

The principal comparison points are listed in Table 1. “v10” denotes a fresh full-capacity v9 baseline. C1 is a light stage-tangent fine-tune. The first attempted C2 was invalid because a tied-multiplier checkpoint was silently loaded into an untied template; an architecture guard was added, and only the corrected C2 should be interpreted scientifically. C10 is the strongest stage-gain variant and the best pre-C16 NaN count. These distinctions explain why an isolated statement such as “C2 failed” can be misleading without a run path and checkpoint choice.

Model	Tanaka NaN	Tanaka div.	Tanaka med.	BF NaN	BF div.	BF med.
v10 baseline	4/32	5/32	0.0123	2/31	2/31	0.00902
C1 stage match	3/32	3/32	0.0151	0/31	0/31	0.00661
Corrected C2 final	3/32	4/32	0.00907	0/31	0/31	0.00735
C10 final	2/32	3/32	0.0138	0/31	0/31	0.00971

Table 1: Primary pre-C16 comparisons under the double-precision harness. The Benjamin–Feir denominator excludes one truth-invalid case.

For a fine-tune driven by an auxiliary rollout or tangent loss, “best validation” and “best physical checkpoint” need not coincide. The ordinary validation score measures the pointwise supervised objective and can select an early epoch before the auxiliary constraint has acted. This was demonstrated by the corrected C2 run: the final checkpoint had the meaningful regularized behavior, whereas the initially evaluated best-validation checkpoint from the misconfigured run was unusable. For C16, final/latest is therefore the primary rollout checkpoint and best-validation is a secondary control.

3.2 Phase I: establish a trustworthy evaluation

Before asking why a model fails, we removed two ways the evaluation itself could lie.

1. A double-precision rollout harness reduced a spurious long-time error floor by factors of roughly 2.5 on Tanaka and 20 or more on several Benjamin–Feir regimes. An order-six oracle displayed the same single-precision floor, proving it was not neural-model error.
2. Reference trajectories with nonfinite states or Hamiltonian drift above the validity threshold were gated out. For example, one Benjamin–Feir initial condition failed in the truth solver while the surrogate remained finite. Counting it against the surrogate would invert the conclusion.

The residual Tanaka NaNs survive both corrections and are therefore genuine surrogate failures.

3.3 Phase II: test the data-coverage explanation

Additional steep-Tanaka data in v9 improved regimes that genuinely had a coverage gap, including the Tanaka g_1 family, but left 4/32 deep-water Tanaka g_0 NaNs. Random and source-stratified scans found no nonfinite base rows. Earlier sign-flip heuristics were not treated as proofs of corruption; the more decisive test came later: the learned DNO remained highly accurate when evaluated directly on the exact trajectories that would eventually fail under rollout.

Two forms of hard-negative training were then explored. Model-harvested pre-cascade states are circular: the inputs already contain the old model’s spectral contamination, and their approximate targets can inherit it. These packs worsened performance. Clean truth-generated perturbations avoid that specific flaw and are valid as local learnability probes, but they constrain a finite list of neighborhoods rather than the universal derivative identity. The clean C13 pack trained successfully, but its evaluation was intentionally deprioritized once the structural diagnosis became available. Thus the conclusion is not “additional data can never help”; it is that ordinary pointwise labels are an inefficient and non-general remedy for an operator derivative defect.

3.4 Phase III: test the integrator explanation

The GL2 stage residual becomes large immediately before nonfiniteness, making time-step failure an attractive primary hypothesis. Substep refinements and halving trials, however, moved the failure time only modestly. A matched 2-by-2 comparison—baseline versus C1, with and without Picard-residual containment—left the NaN counts unchanged in every arm. The containment logic was useful operationally because it isolated bad members of a batch, but it did not rescue the affected trajectories.

This phase did reveal a real training signal. C1’s stage-tangent loss cured one of four baseline Tanaka NaNs and both Benjamin–Feir NaNs. Later C10, targeting the production-filtered $k = 64, \dots, 128$ stage gain, reduced Tanaka NaNs to 2/32. Hence sensitivity control is relevant. The failure to reach zero suggested that the indirect stage proxy was not enforcing the defining differential relation strongly or accurately enough.

3.5 Phase IV: suppress the observed spectrum

Hard learned-block cutoffs, soft high-band penalties, run-time high-band limiters, and a checkpoint-carrying output cap all attacked the visible spectral symptom. None reached zero NaNs with acceptable tails. The strongest soft penalty left its own high-band diagnostic essentially unchanged.

The C10 run-time limiter retained both Tanaka NaNs and worsened the Benjamin–Feir p95. The C14 hard output projection retained 3/32 Tanaka NaNs, produced six divergent Tanaka cases, and raised Tanaka p95 to 1.25.

The cap’s algebra explains its failure. It imposed

$$\|g_{\text{hi}}\| \leq \max(5, 0.1 \|g_{\text{lo}}\|). \quad (7)$$

The dangerous $k = 16, \dots, 32$ content belonged to the unconstrained “low” group, growth of that group increased the permitted high-band norm, and the absolute floor was inactive on some failing states. This is a moving relative envelope, not an absolute stability bound.

Finally, ablations showed why editing a trained checkpoint is hazardous. Block 7 appeared individually amplifying, yet zeroing it changed the Tanaka NaN count from 4/32 to 28/32 and the median error by a factor of 62. Its anti- $|k|$ response was compensating the other blocks. Likewise, removing a large derivative feature from an already trained model breaks learned cancellations and does not tell us whether a clean architecture should expose that feature differently.

Experiment family	Main evidence	Conclusion
Clean v9 retraining	Excellent one-step error, yet the baseline retained 4/32 Tanaka NaNs. Additional steep-Tanaka data cured a separate g_1 coverage failure but not the deep-water g_0 failures.	Pointwise coverage is not the principal remaining defect.
GL2 stage regularization	The best corrected fine-tune reduced Tanaka NaNs from 4/32 to 3/32 and Benjamin–Feir NaNs to zero, but did not reach the zero-NaN target and traded easy-case accuracy.	Directionally useful, but an indirect and incompletely matched constraint.
Step residual checks and halving	Picard residual checks detected terminal noncontractivity. Step refinement and attempted halving moved failure times only slightly; per-sample containment did not change the NaN count.	The integrator exposes the failure after the model state is already contaminated.
Hard block cutoffs and output caps	Hard k -cuts and high-band envelopes reduced selected spectral symptoms but produced poor error tails or retained NaNs. The one-block capped model reached 3/32 Tanaka NaNs and severe p95 regression.	Static culling removes physical response and leaves loopholes in the growing mid-band.
Targeted hard negatives	Model-harvested pre-cascade fine-tunes were harmful; clean truth-targeted packs did not provide a general or journal-explainable mechanism.	Do not substitute case selection for an operator identity.
Exact $G_0 + G_1$	A full 32-case CPU rollout gave 0/32 NaNs and 0/32 divergences, but final η error remained too large (median 0.207, p95 0.410).	The analytic null structure is stable, but needs an accurate higher-order residual.

Table 2: Distilled experimental chronology.

Two broader lessons emerged. First, deleting an apparently dangerous learned block made the model dramatically worse (28/32 rather than 4/32 Tanaka NaNs), showing that the trained eight-block representation relies on fragile inter-block cancellation. Second, increasing or decreasing parameter count did not track rollout stability. The missing ingredient is a structural identity, not capacity alone.

4 Frame-by-frame failure anatomy

Let $z_n^* = (\eta_n^*, \xi_n^*)$ be an exact truth state and z_n the corresponding surrogate-rollout state. Separate the model error into the truth-manifold bias and the response induced by rollout displacement:

$$b_n = \mathcal{G}_\theta(z_n^*) - G(z_n^*), \quad (8)$$

$$i_n = \mathcal{G}_\theta(z_n) - \mathcal{G}_\theta(z_n^*), \quad (9)$$

$$q_n = \frac{\|i_n\|_2}{\|z_n - z_n^*\|_2}. \quad (10)$$

The eventual failures remain accurate on truth states: median one-step relative $G\xi$ errors are 1.87×10^{-4} and 1.72×10^{-4} for representative failing cases 5 and 11. The transverse response, however, grows by four orders of magnitude relative to that bias.

Case	Time	State RMS error	Bias RMS	Induced RMS	Secant gain
5	0.8	1.62×10^{-6}	5.12×10^{-6}	5.56×10^{-6}	3.43
5	58.4	2.08×10^{-3}	5.13×10^{-6}	1.24×10^{-2}	5.95
5	66.4	1.67×10^{-2}	5.12×10^{-6}	1.35×10^{-1}	8.11
11	142.4	3.84×10^{-3}	3.15×10^{-6}	2.71×10^{-2}	7.04
11	150.4	9.02×10^{-3}	3.15×10^{-6}	7.59×10^{-2}	8.42
Healthy 0	200.0	5.91×10^{-5}	1.88×10^{-7}	1.96×10^{-5}	0.33
Healthy 6	200.0	3.76×10^{-4}	1.56×10^{-6}	6.71×10^{-4}	1.78

Table 3: Growth of the model response to its own state error.

Spectrally, the induced response first appears in $k = 32, \dots, 64$, then spreads through $k = 64, \dots, 128$. This precedes the abrupt terminal event. Thus the NaN is not born in one exceptional arithmetic operation; it is the endpoint of a resolved spectral cascade.

The ratio q_n is a secant quantity, not an eigenvalue and not a proof that the full Fréchet derivative has norm 8. It is nevertheless the relevant finite-amplitude response along the error direction actually selected by the rollout. The contrast with healthy cases matters more than its absolute value. In failing case 5 the induced response becomes 2.6×10^4 times the truth-state bias, while a healthy case remains at a gain below one at final time.

The frame order provides a causal discriminator:

1. The pointwise truth-state bias is small and nearly stationary.
2. A small state discrepancy develops through ordinary accumulated error.
3. The learned operator maps that discrepancy into a disproportionately large mid-band response.
4. The response broadens spectrally and steepens the state.
5. Only then do base-grid nonlinear products differ materially from their padded counterparts.
6. Finally, the GL2 fixed-point residual ceases to contract and a nonfinite stage is produced.

This ordering rules out “one bad target row directly produces a NaN” and “the first cause is a single aliased product” as descriptions of the observed cases. It does not claim that all possible future NaNs must follow the same path; the conclusion is scoped to the remaining Tanaka population in the audited checkpoints.

5 Hamiltonian inconsistency is the root defect

For a genuine DNO define

$$B = \frac{G(\eta)\xi + \eta_x \xi_x}{1 + \eta_x^2}, \quad V = \xi_x - B\eta_x. \quad (11)$$

The Hadamard shape derivative satisfies, for every perturbation ζ ,

$$D_\eta G(\eta)[\zeta]\xi + G(\eta)(\zeta B) + \partial_x(\zeta V) = 0. \quad (12)$$

This is an identity of operators applied to ξ , not merely an energy conservation statement along one trajectory. Its importance can be seen by a short variational calculation. Let

$$K_\theta(\eta, \xi) = \frac{1}{2} \langle \xi, \mathcal{G}_\theta(\eta)\xi \rangle, \quad (13)$$

where $\langle f, g \rangle = \int_0^{2\pi} fg \, dx$. Assuming self-adjointness of $\mathcal{G}_\theta(\eta)$ and Eq. (12), periodic integration by parts gives

$$D_\eta K_\theta[\zeta] = \frac{1}{2} \langle \xi, D_\eta \mathcal{G}_\theta(\eta)[\zeta]\xi \rangle \quad (14)$$

$$= -\frac{1}{2} \langle \xi, \mathcal{G}_\theta(\eta)(\zeta B_\theta) \rangle - \frac{1}{2} \langle \xi, \partial_x(\zeta V_\theta) \rangle \quad (15)$$

$$= -\frac{1}{2} \langle \mathcal{G}_\theta(\eta)\xi, \zeta B_\theta \rangle + \frac{1}{2} \langle \xi_x, \zeta V_\theta \rangle \quad (16)$$

$$= \frac{1}{2} \int_0^{2\pi} \zeta (V_\theta \xi_x - B_\theta \mathcal{G}_\theta(\eta)\xi) \, dx. \quad (17)$$

Therefore

$$-\frac{\delta K_\theta}{\delta \eta} = \frac{1}{2} (B_\theta \mathcal{G}_\theta(\eta)\xi - V_\theta \xi_x). \quad (18)$$

Using $V_\theta = \xi_x - B_\theta \eta_x$ and $B_\theta = (\mathcal{G}_\theta \xi + \eta_x \xi_x)/(1 + \eta_x^2)$, the right-hand side becomes

$$-\frac{1}{2} \xi_x^2 + \frac{1}{2} \frac{(\mathcal{G}_\theta(\eta)\xi + \eta_x \xi_x)^2}{1 + \eta_x^2}, \quad (19)$$

which is exactly the nonlinear kinetic contribution used in Eq. (3). This proves the claimed equivalence.

Conversely, tying the learned Fourier multipliers makes $\mathcal{G}_\theta$ self-adjoint and guarantees $\delta K_\theta/\delta \xi = \mathcal{G}_\theta \xi$, but it says nothing about $\delta K_\theta/\delta \eta$. The surface derivative must separately obey Eq. (12). This is the precise sense in which self-adjointness is necessary but not sufficient.

5.1 The measured defect

For diagnosis, the continuum derivative is evaluated by automatic differentiation of the model energy and compared with the geometric expression on the same discrete grid. Both sides undergo the production $|k| \leq 128$ projection. Schematically, the reported relative defect is

$$d_{\text{shape}} = \frac{\|P_{\leq 128}(\xi_{t,\text{geom}} + \delta_\eta K_\theta)\|_{H^1}}{\|P_{\leq 128} \xi_{t,\text{geom}}\|_{H^1} + \varepsilon}. \quad (20)$$

Case/time	Value error	Defect at truth	Defect at rollout	Order-six defect
Healthy 0 / 200.0	6.47×10^{-5}	0.0896	0.2966	2.0×10^{-9}
5 / 33.6	2.34×10^{-4}	1.394	4.975	1.43×10^{-5}
11 / 189.6	2.15×10^{-4}	0.502	1.117	3.43×10^{-6}
22 / 45.6	5.54×10^{-4}	2.724	7.481	3.82×10^{-6}

Table 4: Small pointwise error coexisting with a large shape-derivative defect.

The same normalization, projection, and differentiation pipeline is used for the order-six reference control.

The measured discrepancy, after the production projection to $|k| \leq 128$, is summarized below. The exact order-six DNO passes the same discrete test, ruling out a diagnostic normalization error.

This resolves an apparent contradiction. Gauss–Legendre methods are symplectic for the Hamiltonian vector field supplied to them. The surrogate field is not, in general, the gradient of one Hamiltonian because its η_t and ξ_t components use incompatible derivatives of $\mathcal{G}_\theta$. The time integrator therefore cannot protect the rollout from the learned geometric inconsistency.

One can also phrase the problem through the error equation. Write the truth and surrogate vector fields as F and F_θ , and let $e = z_\theta - z$. Before the error is large,

$$e_t = \underbrace{F_\theta(z) - F(z)}_{\text{truth-manifold bias}} + \underbrace{DF_\theta(z)e}_{\text{propagation by learned tangent map}} + O(\|e\|^2). \quad (21)$$

Supervised one-step fitting directly attacks the first term. The aligned NPZ analysis shows that this term stays small while the second becomes dominant. The Hadamard defect identifies a specific incorrect component of DF_θ , rather than appealing vaguely to “error compounding.” The spectral cascade is error compounding, but it is error compounding through a geometrically wrong tangent operator.

6 The missing spectral null form

The exact first Craig–Sulem correction is

$$G_1(\eta)\xi = -G_0(\eta G_0\xi) - \partial_x(\eta \partial_x \xi). \quad (22)$$

Its Fourier interaction kernel is

$$\widehat{G_1(\eta)\xi}_k = \sum_{k'} \underbrace{(kk' - g_0(k)g_0(k'))}_{\mathcal{K}_1(k,k';h)} \widehat{\eta}_{k-k'} \widehat{\xi}_{k'}, \quad g_0(k) = |k| \tanh(h|k|). \quad (23)$$

The sign is worth checking. The first term in Eq. (22) contributes $-g_0(k)g_0(k')$, while $-\partial_x(\eta \partial_x \xi)$ contributes $+kk'$. In deep water, $g_0(k) = |k|$, so $\mathcal{K}_1 = 0$ whenever k and k' have the same sign. This is not a learned small coefficient: it is exact cancellation between two individually large derivative terms. At finite depth the cancellation is exponentially accurate when $h|k|$ and $h|k'|$ are large, as they are near the observed modes.

The null form is physically important in the failing Tanaka regime. A transfer-kernel probe at a truth surface found that the learned model maps $k_{\text{in}} = 117$ to $k_{\text{out}} = 115$ with amplitude 985.67, whereas the order-six reference amplitude is 1.47×10^{-5} , a ratio of approximately 6.7×10^7 . The flat diagonal response at this mode is about 5.9×10^4 , so 985.67 is only about 1.7% of the diagonal

amplitude. That explains why an aggregate pointwise error metric can appear excellent while a repeatedly injected off-diagonal sideband matters dynamically. The learned model can represent the cancellation but does not enforce it.

Across all sampled Tanaka truth states, replacing G_0 by exact $G_0 + G_1$ reduces median relative $G\xi$ error from 0.1496 to 0.00555 and p95 from 0.2089 to 0.01373. More importantly, it makes the high-frequency null form exact rather than dependent on cancellation among learned blocks.

The analytic-only rollout is a particularly discriminating experiment. Exact $G_0 + G_1$, with no cap and no learned residual, yields 0/32 NaNs and 0/32 divergences. It nevertheless gives final η median 0.2071 and p95 0.4096. Hence two statements are simultaneously true:

1. the first-order null structure removes the catastrophic mechanism in this test; and
2. neglected $G_2, G_3, \dots$ terms are essential for useful long-time phase and amplitude accuracy.

This is why the proposed solution is not to abandon learning, but to prevent the learned residual from overwriting the analytically known first-order term.

7 Aliasing and the role of the integrator

For truncated spectra, a physical-space product computes a circular convolution. If f and g contain modes near the retained cutoff, terms with true wavenumber outside the grid wrap into lower modes. The nonlinear Zakharov right-hand side contains products and a rational expression in η_x , ξ_x , and $G\xi$, so this mechanism becomes severe once the surrogate has populated its upper retained bands.

Xu and Guyenne explicitly identify aliasing in both the DNO expansion and the equations of motion, and extend the spectra by a factor of two for nonlinear terms in the equations of motion [1, Sec. 3.3]. Our order-six reference DNO already uses padding, but the original surrogate and truth rollout paths evaluated the quadratic Zakharov products at base resolution. The correction now implemented is:

1. Fourier-interpolate η_x , ξ_x , and $G\xi$ onto a grid with twice as many points;
2. evaluate the nonlinear expression pointwise on that padded grid; and
3. truncate the result back to the production spectrum before applying the usual retained-mode projection.

Two-times Fourier padding is therefore a necessary discretization correction. On healthy trajectories, base-grid and padded right-hand sides agree to about 10^{-11} . Near the last finite frames, the discrepancy becomes large:

Case	Last finite time	Relative RHS difference	First $> 10^{-3}$	First > 0.1
Healthy 0	200.0	2.71×10^{-11}	never	never
Healthy 6	200.0	1.78×10^{-11}	never	never
5	33.6	0.753	32.0	33.6
11	189.6	0.819	187.2	188.8
22	45.6	0.0829	45.6	never

Table 5: Difference between base-grid and two-times-padded Zakharov right-hand sides.

The large terminal discrepancy alone does not establish that aliasing started the cascade. A causal replay does: restarting each failing state with padded products delays the first nonfinite

value by only 0.05, 0.08, and 0.05 time units for cases 5, 11, and 22. A healthy replay agrees to 3.5×10^{-15} . Aliasing is therefore a selective terminal amplifier, not the initiator in these cases. The padding fix should be retained regardless of C16’s outcome because it restores the intended pseudospectral discretization without perturbing resolved healthy trajectories.

Likewise, more Picard work or smaller steps cannot undo the preceding operator-induced drift. A residual check remains valuable as a reported fail-safe, and adaptive halving may rescue milder future failures, but it is not a substitute for restoring Hamiltonian compatibility.

There is no conflict with the well-known long-time advantages of symplectic integration. Those advantages concern the discretization of a Hamiltonian vector field and usually imply bounded, structured energy error over long times, rather than exact conservation at each step. Here the learned vector field fails the compatibility calculation in Eq. (17). Moreover, the practical scheme uses a fixed number of nonlinear iterations, so once the state enters a high-gain region it may not even realize the exact implicit GL2 map. The integrator is important, but it is downstream of the operator geometry.

8 C16: the resulting model and objective

The proposed clean retraining uses

$$\mathcal{G}_\theta(\eta) = G_0 + G_1(\eta) + \mathcal{R}_\theta(\eta), \quad \mathcal{R}_\theta(\eta) = O(\eta^2). \quad (24)$$

The analytic G_1 path is uncut and evaluated with double-precision spectral arithmetic. The learned residual uses eight tied, self-adjoint blocks with 256 branches each. A bias-free feature trunk followed by $f \mapsto f \tanh(f)$ ensures

$$\mathcal{R}_\theta(0)\xi = 0, \quad D_\eta \mathcal{R}_\theta(0)[\zeta]\xi = 0 \quad (25)$$

for every parameter value. The residual sum is scaled by $1/\sqrt{8 \cdot 256}$, making the size of a fresh Adam update independent of the chosen parallel rank without changing the function class.

8.1 Residual parameterization and its invariants

Suppressing depth and normalization factors, the residual has the schematic form

$$\mathcal{R}_\theta(\eta)\xi = \frac{1}{\sqrt{N_b N_r}} \sum_{b=1}^{N_b} \sum_{r=1}^{N_r} M_{b,r}(D, h) [\Phi_{b,r}(\eta) M_{b,r}(D, h)\xi], \quad N_b = 8, \quad N_r = 256. \quad (26)$$

The same real Fourier multiplier appears on both sides of the pointwise multiplication. For fixed η , each term is self-adjoint:

$$\langle \psi, M[\Phi M\xi] \rangle = \langle M\psi, \Phi M\xi \rangle = \langle \Phi M\psi, M\xi \rangle = \langle M[\Phi M\psi], \xi \rangle. \quad (27)$$

The proof uses self-adjointness of the real multiplier and of multiplication by the real function Φ . Summation and positive scalar normalization preserve the property.

For the order statement, let the bias-free eta trunk produce $f(\eta) = L\eta + O(\|\eta\|^2)$ near the flat surface. Since $\tanh f = f + O(f^3)$,

$$f(\eta) \tanh f(\eta) = (L\eta)^2 + O(\|\eta\|^3). \quad (28)$$

All subsequent eta projections are bias-free. Consequently every $\Phi_{b,r}(0) = 0$ and $D_\eta \Phi_{b,r}(0) = 0$, proving the two structural zeros. The bounded-growth lift was chosen instead of raw f^2 : it retains

quadratic behavior near zero but grows only linearly in $|f|$ when derivative features are large. Learned mixing before the lift still permits general quadratic cross-interactions; this is not a restriction to literal squares of the physical input channels.

The architecture contains 860,928 trainable parameters. It retains the full eight-block capacity because the one-block and block-deletion experiments did not support a capacity reduction. There is no hard spectral cutoff, no output cap, no PSD hinge, and no Tanaka-specific data pack in C16. Thus a successful result can be attributed to the null-structured backbone and direct shape constraint rather than to spectral culling.

8.2 Finite-secant Hadamard loss

The randomized finite-secant residual is

$$S_\theta = \frac{\mathcal{G}_\theta(\eta + \epsilon\zeta)\xi - \mathcal{G}_\theta(\eta)\xi}{\epsilon} + \mathcal{G}_\theta(\eta)(\zeta B_\theta) + \partial_x(\zeta V_\theta), \quad (29)$$

and the training objective is

$$\mathcal{L} = \mathcal{L}_{\text{data}} + \lambda_{\text{shape}} \mathbb{E}_\zeta \frac{\|WP_{\leq 128} S_\theta\|_2^2}{\|WP_{\leq 128} (\mathcal{G}_\theta(\eta)(\zeta B_\theta) + \partial_x(\zeta V_\theta))\|_2^2 + \delta}. \quad (30)$$

Here W is the supervised Sobolev weight. Relative, zero-mean, inverse-Sobolev probes cover low and mid/high modes without selecting specific Tanaka cases. The regularizer is evaluated in double precision on a global microbatch of eight every 16 optimizer steps, with $\lambda_{\text{shape}} = 10^{-2}$ introduced through a warmup.

More explicitly, a random Fourier probe is projected to zero mean and scaled by an inverse Sobolev symbol so its high modes do not dominate merely because derivatives multiply by $|k|$. Its physical RMS is chosen relative to the RMS of the current η , subject to a small floor, and $\epsilon \in [10^{-3}, 3 \times 10^{-3}]$ is sampled as a relative secant size. The two model evaluations use exactly the production input normalization, output de-normalization, filtering, and output-mean subtraction. Both the secant residual and its normalizing forcing are measured after projection to $|k| \leq 128$ in the same Sobolev metric used by supervised training.

The denominator in Eq. (30) is essential. An absolute residual would overweight large-amplitude samples and permit the loss scale to change with probe amplitude. The relative form asks whether the identity error is small compared with the physical terms that should cancel. The floor δ only regularizes nearly trivial forcing. A finite secant, rather than an exact Jacobian norm, avoids differentiating through a full Jacobian and keeps the cost compatible with sparse microbatch evaluation.

This loss is related to, but cleaner than, the earlier GL2 stage regularizer. The stage loss asks whether one particular discretized short-time map behaves well under selected perturbations; it mixes DNO error, time-discretization details, filters, and stage construction. The Hadamard loss asks directly for the continuum operator identity that makes the Zakharov vector field Hamiltonian. It is therefore cheaper to interpret and easier to defend in a mathematical presentation.

8.3 Optimizer calibration

The first implementation exposed a separate optimizer parameterization issue: hundreds of thousands of zero-initialized projection weights received aligned Adam sign updates and were summed without rank normalization. Before the normalization, a single real-state step at learning rate 2×10^{-5} changed loss from 0.02463 to 0.3201. After normalization, the same step reduced it to 0.02174. A

1% v9 GPU smoke then completed with train loss 0.01361 and validation loss 0.01311, with no surge or nonfinite metric.

The mechanism of the initial optimizer instability deserves emphasis because it is independent of the NaN-rollout diagnosis. At initialization, the analytic baseline already has small loss and the per-block eta-to-branch projection kernels are zero. Adam normalizes each nonzero coordinate by its own second moment, so the first update is approximately a sign step in each of 524,288 projection coordinates. All coordinates are trained against the same residual and their function-space effects add coherently across eight blocks and 256 branches. A parameter-space learning rate that was harmless in a warm-started model therefore produced an enormous first function-space step in the fresh model.

The factor $1/\sqrt{N_b N_r} = 1/\sqrt{2048}$ is a fan-in normalization for the parallel residual paths. It preserves exact $G_0 + G_1$ at initialization, the order-two zero, self-adjointness, parameter count, and the set of functions representable after rescaling downstream weights. CPU calibration gave

Parameterization	Loss before step	Loss after AdamW step at 2×10^{-5}
Unnormalized order-two residual	0.024634	0.320115
Fan-in-normalized residual	0.024634	0.021744

The data-gradient to Hadamard-gradient norm ratio was 0.0662, and the chosen Hadamard weight 0.01 makes the auxiliary gradient influential but not dominant at initialization. Both the learning rate and Hadamard weight are warmed up over 500 steps.

8.4 What C16 guarantees, and what it does not

For every value of its trained parameters, C16 guarantees:

- exact $G_0 + G_1$ behavior through first order at $\eta = 0$;
- linearity and self-adjointness in ξ for the learned residual;
- no learned $O(1)$ or $O(\eta)$ term that can overwrite the analytic first-order null form; and
- a consistently two-times-dealiased Zakharov nonlinear right-hand side.

It does *not* guarantee the full nonlinear Hadamard identity for finite η ; that is encouraged statistically by Eq. (30). Nor does it prove bounded rollouts, exact energy conservation, or uniform operator error outside the sampled distribution. Calling it an “explicit bounded residual architecture” would therefore be inaccurate. Its principle is geometric compatibility and asymptotic null structure, not a hard global norm bound.

9 Current experiment and acceptance criteria

At the time of writing, a budget-capped C16 retraining is running on two H100 GPUs in the personal Modal workspace. It uses clean v9 data, batch size 1024, learning rate 2×10^{-5} , 500-step learning-rate and Hadamard warmups, and 16 epochs. The abbreviated run is intended to obtain the strongest checkpoint possible within approximately 30 dollars of workspace credit; it is not assumed that epoch 16 is fully converged.

The first two complete epochs are healthy:

Epoch	Train loss	Validation loss	Hadamard loss	Defect RMS
1	0.009919	0.008740	0.04381	0.1700
2	0.008472	0.008376	0.03104	0.1482

The regularizer was active on 373 of 5964 batches in each epoch, matching the intended frequency 1/16. Its mean effective weight reached 9.57×10^{-3} in epoch 1 and the full 10^{-2} in epoch 2. These are encouraging training diagnostics, not rollout evidence. In particular, the random-probe defect RMS cannot yet be compared directly with the case-specific defect table without running the same held-state diagnostic.

An earlier Modal attempt stopped one third of the way through epoch 1 because the local monitoring client was explicitly interrupted; Modal logs show a cancellation signal rather than a Python, JAX, or model exception. No checkpoint was produced by that attempt. The current 16-epoch run is a fresh relaunch and should be judged independently.

The structural hypothesis should be accepted only if rollout evaluation shows all of the following:

1. zero Tanaka NaNs and divergences among 32 truth-valid cases;
2. zero scored Benjamin–Feir NaNs and divergences among 31 cases;
3. no material median or p95 regression relative to the strongest prior checkpoints;
4. at least a two-order-of-magnitude reduction in the failing-state shape defect;
5. no pre-cascade secant-gain growth toward 7–8; and
6. healthy GL2 fixed-point residuals without relying on output caps, hard spectral cutoffs, or case-specific damping.

The main scientific conclusion does not depend on the outcome of one training run: long-rollout stability requires approximating the geometry of the DNO in addition to its values. C16 is the first experiment in this sequence that directly enforces the missing first-order null form and shape derivative while retaining the capacity needed for higher-order accuracy.

9.1 Planned post-training decision sequence

To avoid another ambiguous experiment, evaluation should proceed in the following order:

1. Verify checkpoint/config integrity and run the held-state Hadamard defect table on the final checkpoint. This is cheap and directly tests whether the intended loss changed the diagnosed quantity.
2. Evaluate the final checkpoint on Tanaka g_0 and Benjamin–Feir g_1 with cached truth, the double-precision harness, batched surrogate execution, and no run-time limiter or damping.
3. If final is negative, evaluate best-validation only as a checkpoint selection control; do not reinterpret it as a separate architecture.
4. For every remaining failure, regenerate the frame-aligned secant, spectral-band, RHS-alias, and GL2-residual traces. The question is not merely whether the failure count moved, but whether the causal signature changed.
5. Only after the two decisive regimes pass should the full multi-regime suite be run to rule out regressions in shallow water, random seas, Stokes waves, and modal instability.

10 Alternative explanations, objections, and falsifiers

This section states the questions most likely to arise in a rigorous oral discussion and the limits of the current answers.

10.1 Could the data still contain subtle corruption?

Yes, in the weak logical sense that no finite audit proves every stored number correct. What has been ruled out is the form of corruption needed to explain the observed failure mechanism. Random and source-stratified scans found no nonfinite base rows; sign and spectral heuristics did not reveal a population aligned with the failing cases; and, most importantly, the learned DNO has small value error when evaluated on the exact failing truth states. A corrupted label population severe enough to seed the reported sideband directly should normally appear as increased error on those states. Instead, the large error appears when the input is moved in the rollout’s transverse direction.

A useful remaining audit would compare independently regenerated order-six labels against stored v9 labels on a stratified sample, including source, depth, steepness, and high-band-energy tails. Such an audit would strengthen dataset provenance but would not replace the already measured shape identity.

10.2 Does small value error plus large derivative error necessarily cause a NaN?

No. A large local derivative defect is a mechanism and risk factor, not by itself a theorem of finite-time blow-up. The causal case is empirical: the defect is present at the failing states, the realized secant direction grows, its spectrum matches the subsequent cascade, and analytic restoration of the first-order null form removes all NaNs in the 32-case test. A formal theorem would require stability bounds for the coupled semi-discrete system and a uniform neighborhood estimate for $D\mathcal{G}_\theta$, neither of which is currently available.

10.3 Could GL2 still be the primary issue?

The fixed number of Picard iterations is certainly where nonfiniteness is finally emitted, and a more robust nonlinear solver could extend some trajectories. It is unlikely to be the primary issue in the audited cases because the model-induced mid-band response is already orders of magnitude above its truth-state bias before the stage residual fails; substep refinement and padded-RHS replay delay failure only slightly; and healthy cases use the same integrator without trouble. A definitive follow-up could solve selected near-terminal stages with Newton–Krylov to high tolerance. If the converged implicit map returned to a healthy spectral state, the present conclusion would need revision. If it advanced the already contaminated state, it would confirm the current ordering.

10.4 Why not learn the whole Hamiltonian and differentiate it?

That is a principled alternative. A scalar Hamiltonian network with $\eta_t = \delta_\xi H_\theta$ and $\xi_t = -\delta_\eta H_\theta$ would be Hamiltonian by construction. The tradeoff is computational: rollout requires derivatives of a nonlocal learned energy, training DNO values can involve mixed derivatives, and encoding the known linear dispersion and Craig–Sulem structure is less direct. C16 keeps the efficient DNO interface while enforcing the identity that makes that interface Hamiltonian. If C16 fails despite a substantially reduced measured shape defect, a direct energy architecture becomes the most principled next architectural family.

10.5 Why not impose a global Lipschitz or spectral norm bound?

A uniform bound could control amplification but would not distinguish physical from unphysical high-frequency response. The DNO is an order-one nonlocal operator; its norm naturally scales with wavenumber, and nonlinear wave dynamics legitimately transfer energy across modes. A naive global Lipschitz constant or hard cutoff risks destroying exactly the dynamics we are trying to approximate. The Hadamard identity is selective: it constrains how the large terms must cancel, rather than requiring every term to be small.

10.6 What result would falsify the C16 mechanism?

The strongest falsifier would be: the final C16 checkpoint reduces the held shape defect by two orders of magnitude and preserves the same-sign G_1 null form, yet the same cases still develop the same pre-cascade secant gain and spectral signature. That would show that the diagnosed identity is not the controlling mechanism. Other negative outcomes are less decisive:

- If the shape defect does not improve, the loss estimator, weight, or 16-epoch budget failed to enforce the hypothesis; the hypothesis itself was not tested strongly.
- If the defect improves but one-step accuracy is poor, the structural constraint may be correct but the residual architecture or optimization is inadequate.
- If NaNs disappear but p95 error becomes unacceptable, C16 is a stability result but not a useful surrogate.

A Chronological experiment map

This appendix is deliberately redundant with `EXPERIMENTS.md`. Its purpose is to make the logic reconstructible without memorizing labels.

Handle	Intervention	Outcome	What it taught us
v3–v10	Full-feature CS-DNO baselines; progressively broader clean data and a double-precision rollout harness.	Much better one-step and median rollout accuracy than generic FNOs, but persistent rare Tanaka NaNs (v10: 4/32).	The Craig–Sulem inductive bias is valuable, but pointwise accuracy and broad data coverage do not ensure rollout stability.
C1	Light GL2 stage-tangent matching fine-tune.	Tanaka 4/32 to 3/32 NaNs; Benjamin–Feir 2/31 to 0/31.	Off-manifold sensitivity is a real lever, though the indirect proxy is incomplete.
C2, first launch	Stage match plus gain, but resumed a tied checkpoint into an untied architecture.	32/32 NaNs in both decisive regimes.	Invalid experiment; motivated strict resume architecture and parameter-count guards.

Handle	Intervention	Outcome	What it taught us
C2, corrected	Same intended loss with tied architecture preserved; final checkpoint evaluated after the auxiliary loss acted.	Tanaka 3/32 NaNs, median 0.00907; BF 0/31, median 0.00735.	Final versus best checkpoint matters for auxiliary-loss fine-tunes; still short of zero NaNs.
GL2 guard	Picard residual detection, attempted halving, then per-sample masking.	Detected but did not reduce NaN counts; early fleet-wide logic could couple one bad sample to the whole batch.	GL2 noncontractivity is downstream; batch-local safety logic must remain per sample.
Block-7 ablation	Zero an apparently amplifying learned block.	Tanaka NaNs rose from 4/32 to 28/32; median error rose 62-fold.	Trained blocks form a fragile compensating sum; post hoc deletion is not an architectural study.
v11-D	80-fold smaller model plus cutoff, stage regularizer, and PSD hinge.	Same Tanaka NaN count as baseline and 12–20-fold median regressions.	A kitchen-sink run is hard to interpret and reduced capacity was not a cure.
v11-E	Full capacity with hard learned-block k -cut isolated.	Did not establish a useful zero-NaN/accuracy tradeoff.	Brute-force band removal is not the preferred mechanism.
C6/C7	Model-harvested pre-cascade hard negatives.	Harmful or non-general fine-tunes.	Contaminated off-manifold packs can feed the model’s own failure back into training.
C8/C10	Increasingly matched stage-gain regularization, with C10 using the production filter and $k = 64, \dots, 128$ target band.	C10 reached the best pre-C16 count, 2/32 Tanaka and 0/31 BF NaNs, but regressed some medians/tails.	The diagnosed spectral band matters; indirect regularization still does not restore the defining operator identity.
C9	Stronger soft full-output high-band envelope penalty.	Penalty diagnostic barely moved; Tanaka remained 3/32 NaNs and accuracy regressed.	Simply increasing this penalty weight is not an effective axis.
C10 + runtime limiter	Cap learned $G\xi$ high-band output during rollout.	Tanaka stayed 2/32 NaNs; BF p95 worsened.	The operational guard does not reach zero and can damage tails.
C11/C13	Clean truth-generated hard-negative packs near selected remaining failures.	Training was feasible; C13 evaluation was superseded before a final claim.	Useful as a local learnability probe, not a universal or satisfying paper-level mechanism.
C12	One-step pushforward/filtering fallback prepared from C2.	Not promoted after the stronger structural diagnosis.	Rollout-state training remains an option, but it mixes integrator and operator effects.

Handle	Intervention	Outcome	What it taught us
C14	One-block, checkpoint-carrying full-output cap.	Best finite checkpoint: 3/32 Tanaka NaNs, 6/32 divergences, median 0.162, p95 1.253; training later became nonfinite.	The cap has a mid-band/floor loophole and severe accuracy cost; one-block capacity was also a confounder.
C15	Eight-block capacity control for the C14 cap.	Interrupted before epoch 1 and before any checkpoint.	Inconclusive as a capacity ablation, but not worth relaunching because the cap mechanism itself was analytically rejected.
Root-cause audit	Frame alignment, realized secants, automatic-differentiation shape check, transfer kernel, analytic $G_0 + G_1$, and padded-RHS replay.	Small truth-state value error coexists with $O(1)$ shape defect; exact $G_0 + G_1$ is stable; RHS padding delays rather than cures failure.	The DNO shape derivative/Hamiltonian compatibility is the primary missing constraint.
C16	Fresh exact $G_0 + G_1$, full self-adjoint $O(\eta^2)$ residual, randomized Hadamard loss, and padded RHS.	Structural tests and 1% smoke pass; 16-epoch budget-capped training is in progress.	First clean direct test of the diagnosis.

B Implementation and correctness checks

The implementation was not accepted on training loss alone. The following independent checks passed before the full launch:

- five architecture tests: legacy order-one identity, zero order-two value and first variation, quadratic near-zero scaling, self-adjointness, and isolation of double precision to analytic G_1 ;
- seven Hadamard-loss tests: projected Sobolev energy, probe scaling and bandwidth, exact flat-surface identity, detection of missing G_1 , production normalization/mean behavior, JIT and mixed-precision gradients, and invalid-configuration handling;
- two dedicated padded-RHS tests: equality on bandlimited products and removal of high-frequency quadratic aliasing;
- fourteen existing stage-regularizer tests after routing all nonlinear right-hand sides through the same padded helper; and
- a 59-step two-GPU smoke with finite checkpoints and validation.

Reproducibility record

The detailed chronological record is `EXPERIMENTS.md`. Frame-level statistics and paths to the aligned NPZ archives are recorded in `notes/nan_root_cause_investigation_20260709.md`. The

C16 launcher is `scripts/launch_c16_modal_h100x2.sh`; model, Hadamard-loss, and dealiased-RHS invariant tests are stored alongside their implementations.

References

- [1] L. Xu and P. Guyenne, “Numerical simulation of three-dimensional nonlinear water waves,” *Journal of Computational Physics*, vol. 228, pp. 8446–8466, 2009, doi:10.1016/j.jcp.2009.08.015.