Single-rollout runtime vs error

Blackwell · Modal — RTX PRO 6000

Classical
Gradient EMA — saved accuracy
Neural — small (saved accuracy)
Neural — full (saved accuracy)
M6 time
New candidate optimized network
Spectral DNO — weight EMA (PyTorch)

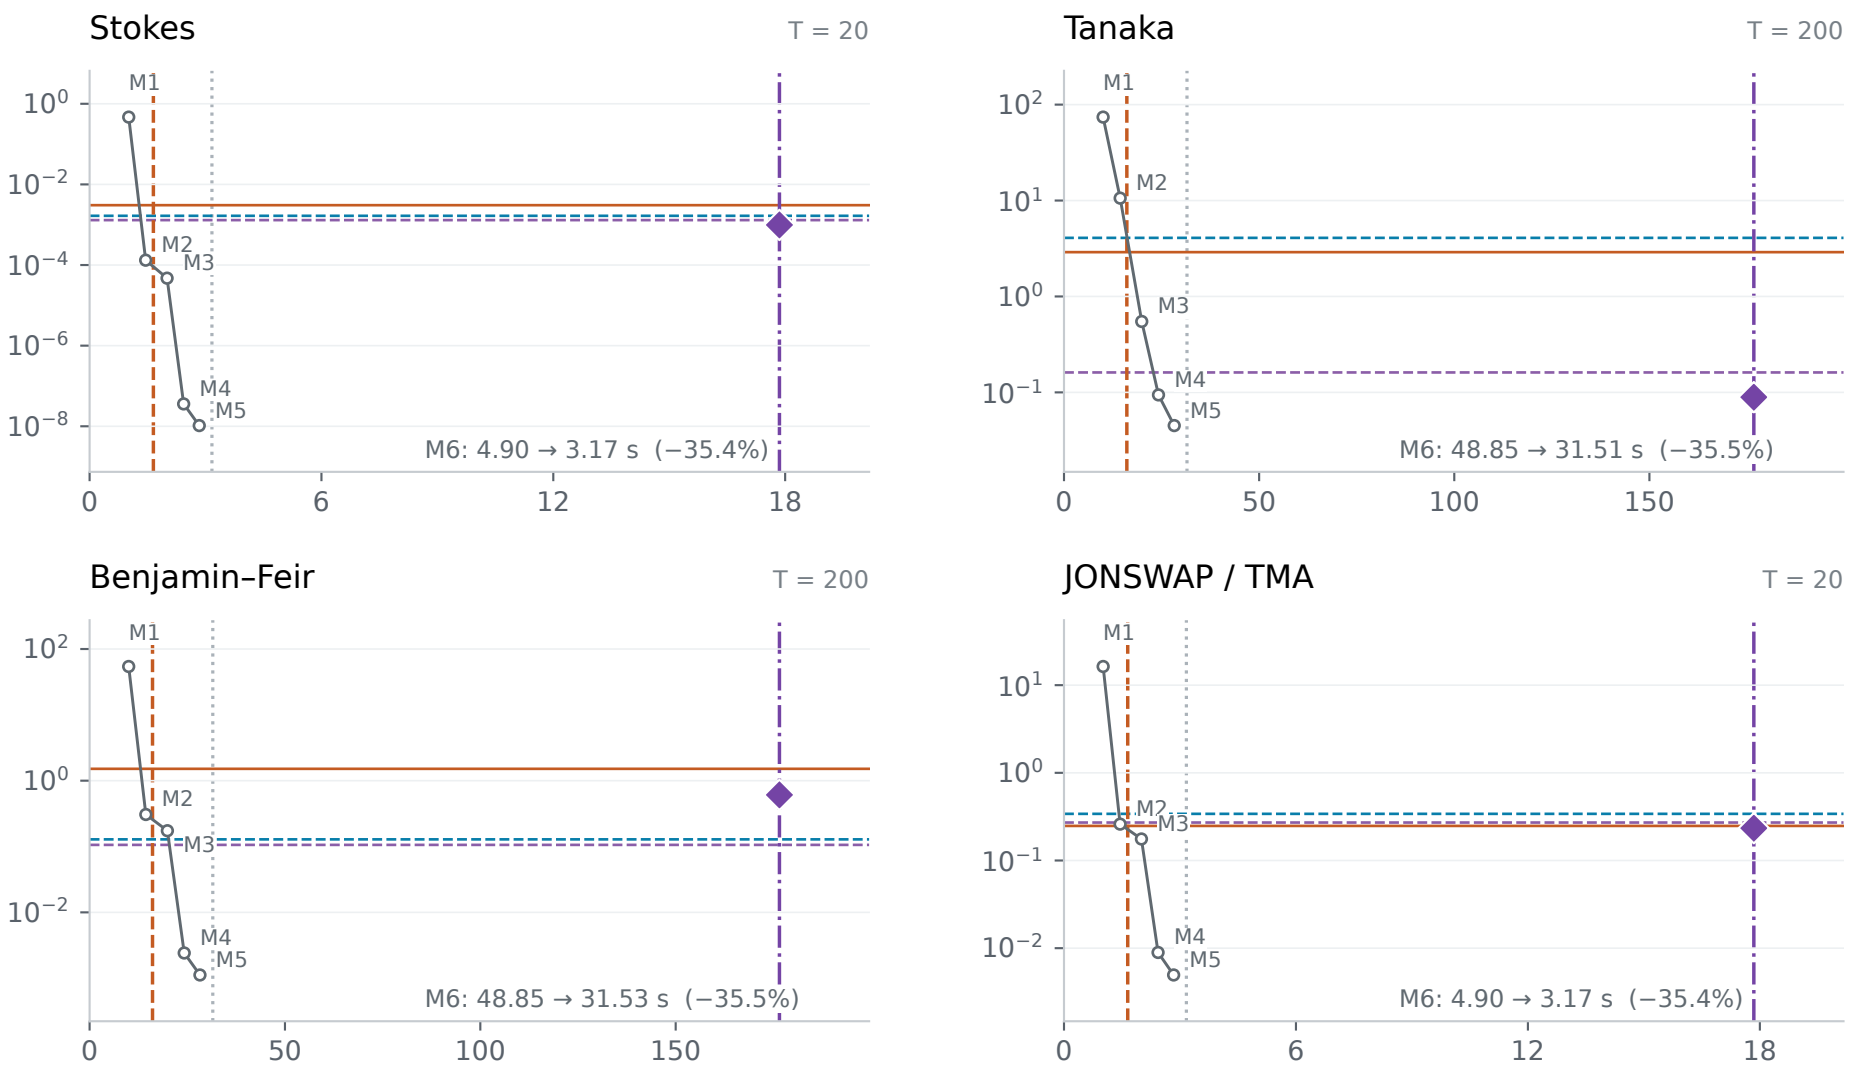

Ada · Local — RTX 6000

Classical
Gradient EMA — saved accuracy
Neural — small
Neural — full
M6 time
New candidate optimized network
Spectral DNO — weight EMA (runtime only)

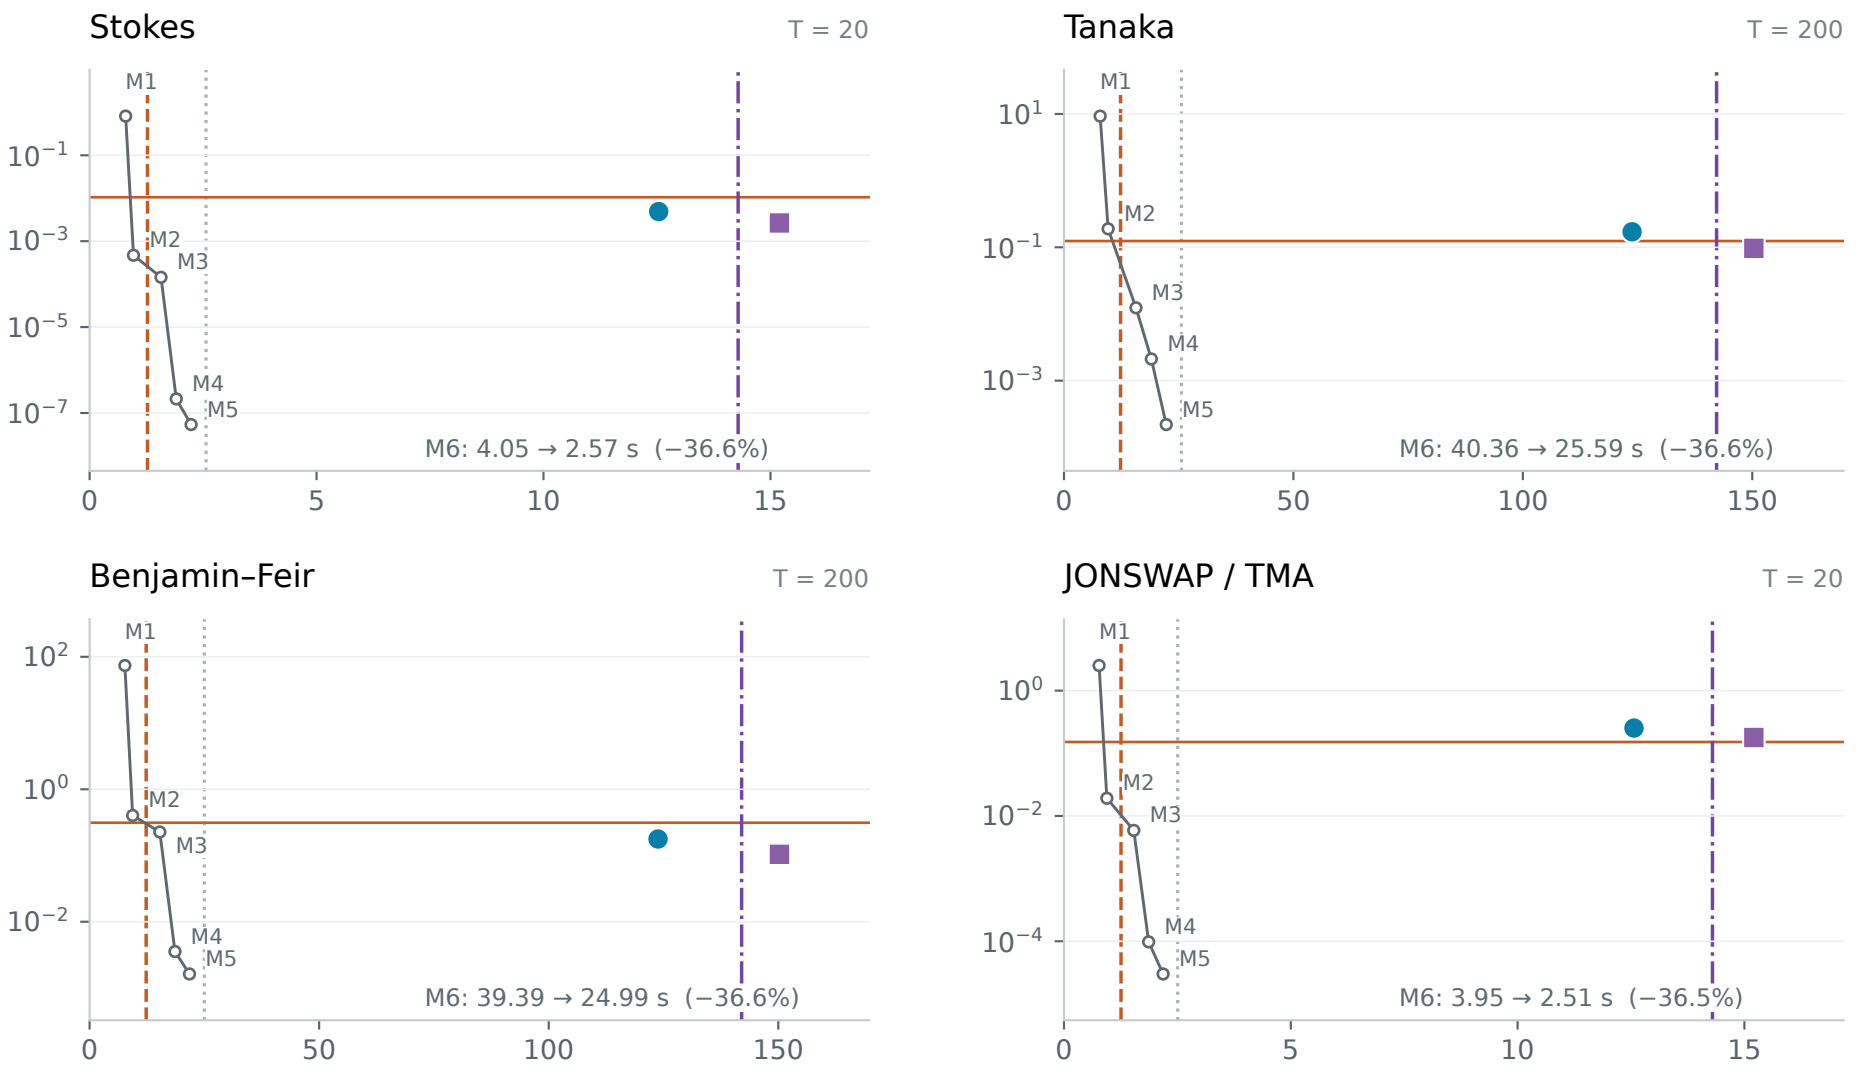

A100 · Modal — A100-SXM4-80GB

Classical
Gradient EMA — saved accuracy
Neural — small (saved accuracy)
Neural — full (saved accuracy)
M6 time
New candidate optimized network
Spectral DNO — weight EMA (PyTorch)

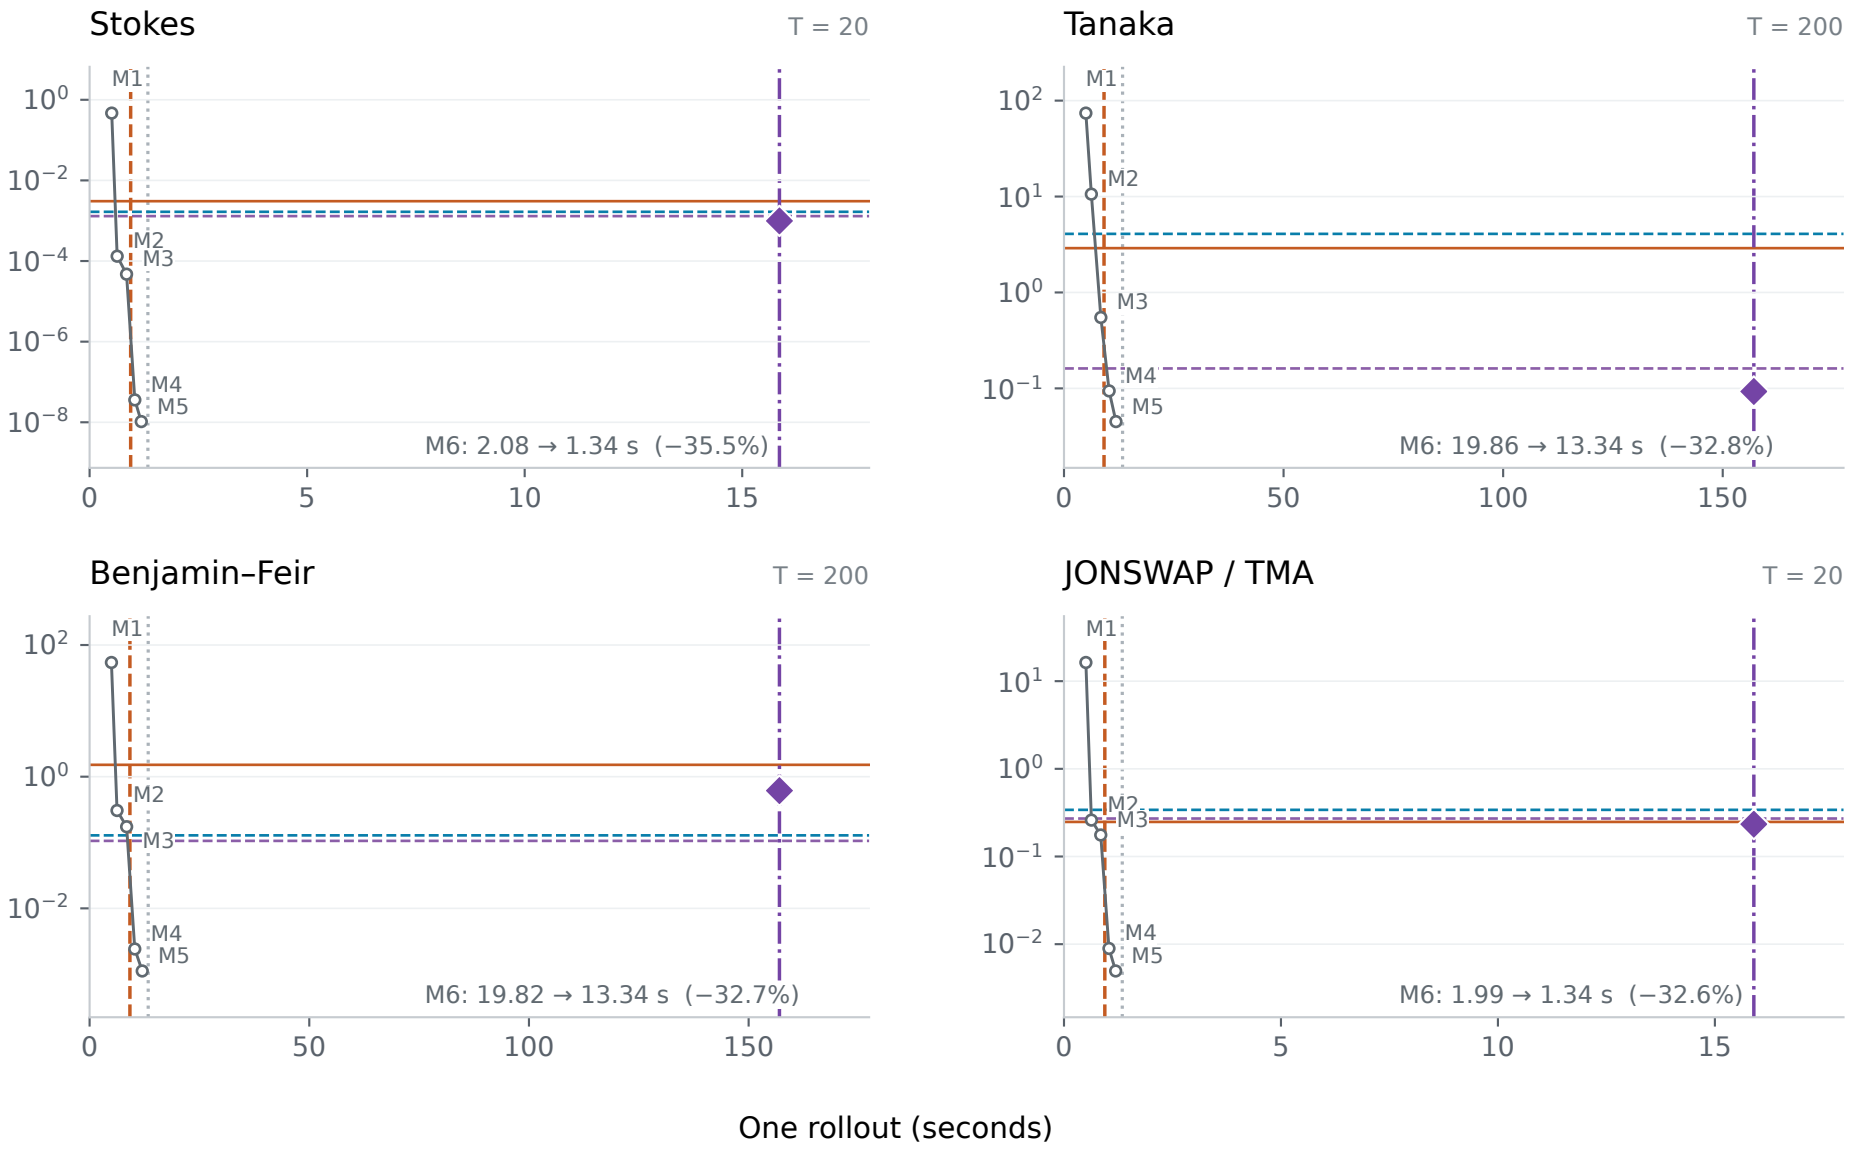

Final surface error (%)

One rollout (seconds)