Single-rollout runtime vs error

Blackwell · Modal — RTX PRO 6000

Classical
Neural — small (saved accuracy)
Neural — full (saved accuracy)
M6 — interleaved FFTs
New candidate optimized network
Spectral DNO — weight EMA (PyTorch)

Stokes — T = 20

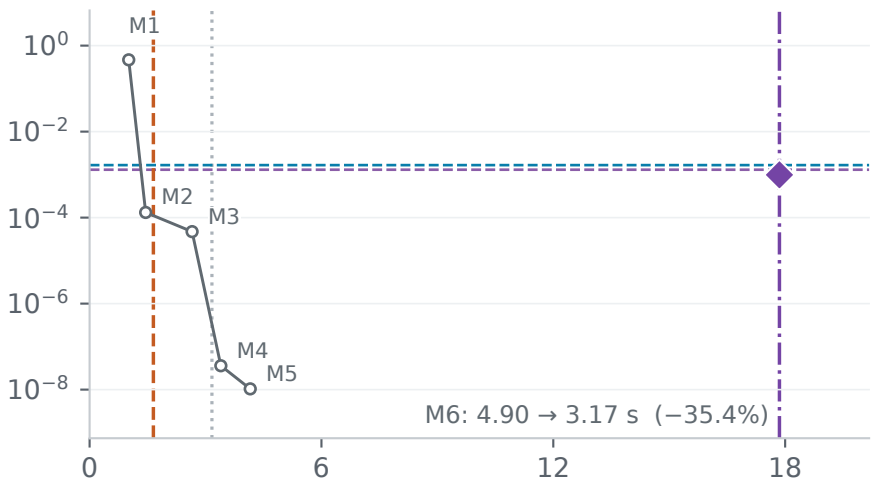

Tanaka — T = 200

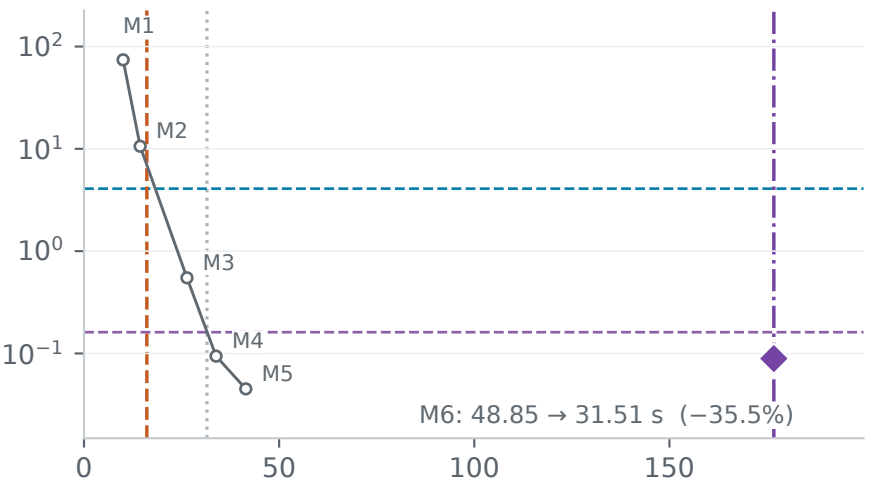

Benjamin–Feir — T = 200

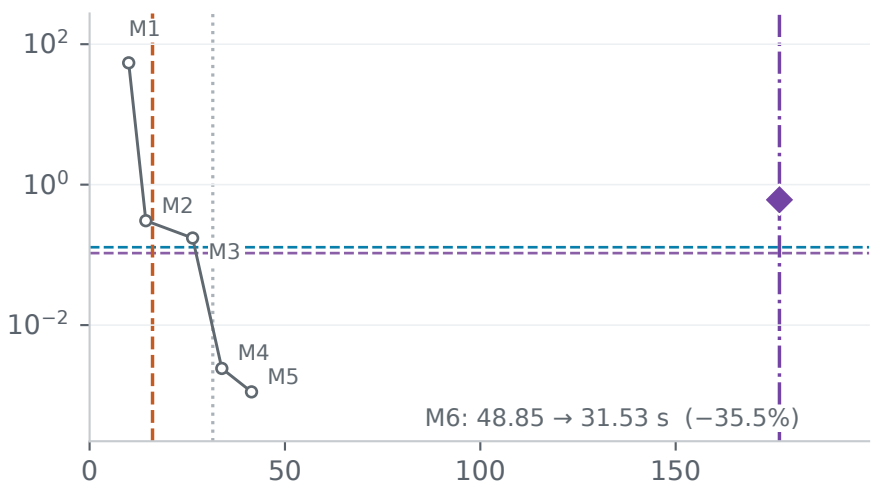

JONSWAP / TMA — T = 20

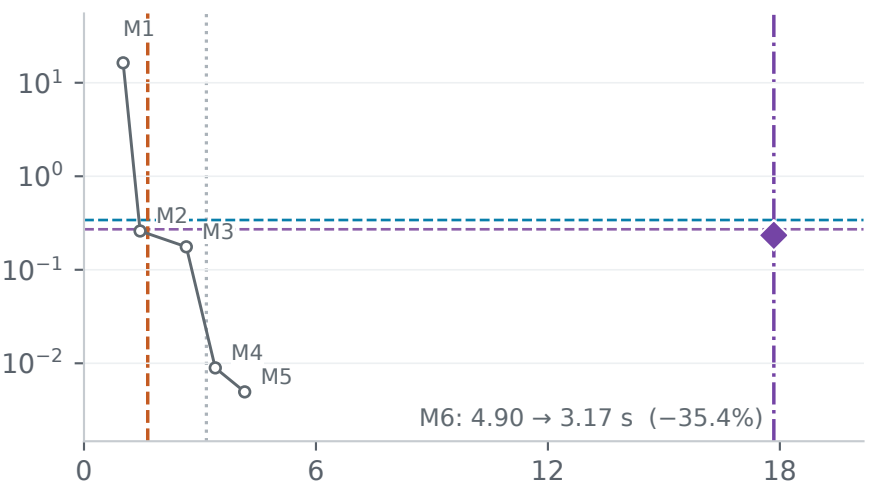

Final surface error (%)

One rollout (seconds)

Ada · Local — RTX 6000

Classical
Neural — small
Neural — full
M6 time
New candidate optimized network
Spectral DNO — weight EMA (runtime only)

Stokes — T = 20

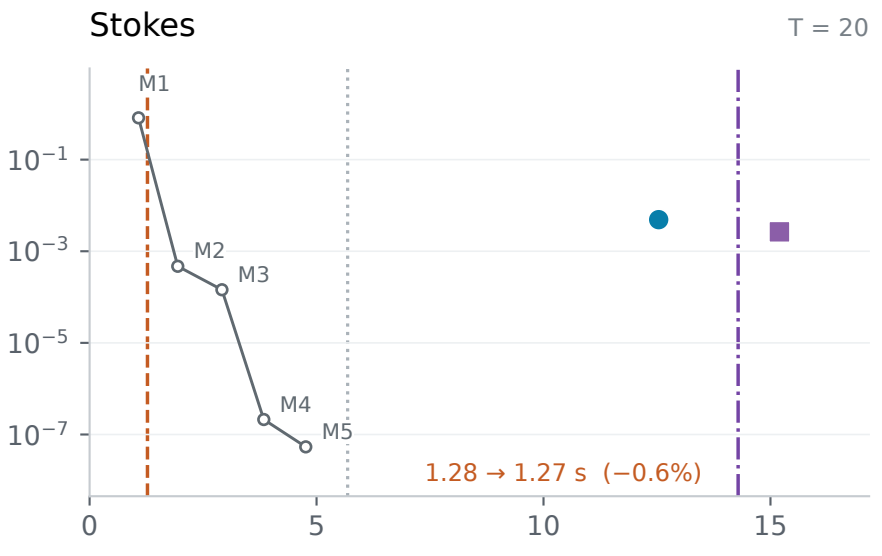

Tanaka — T = 200

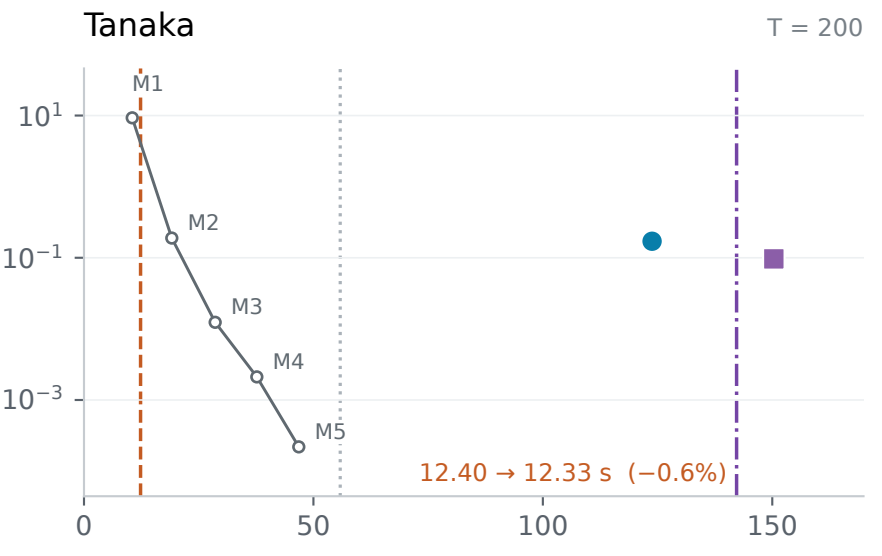

Benjamin–Feir — T = 200

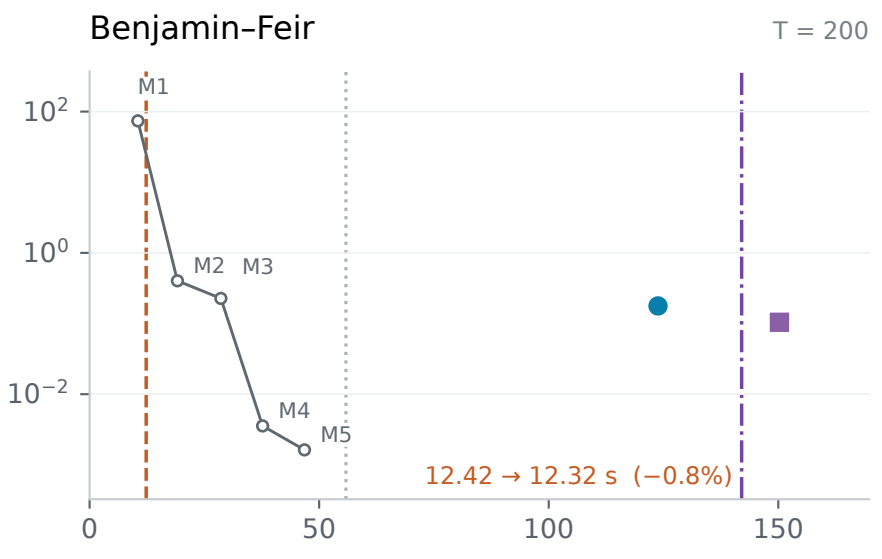

JONSWAP / TMA — T = 20

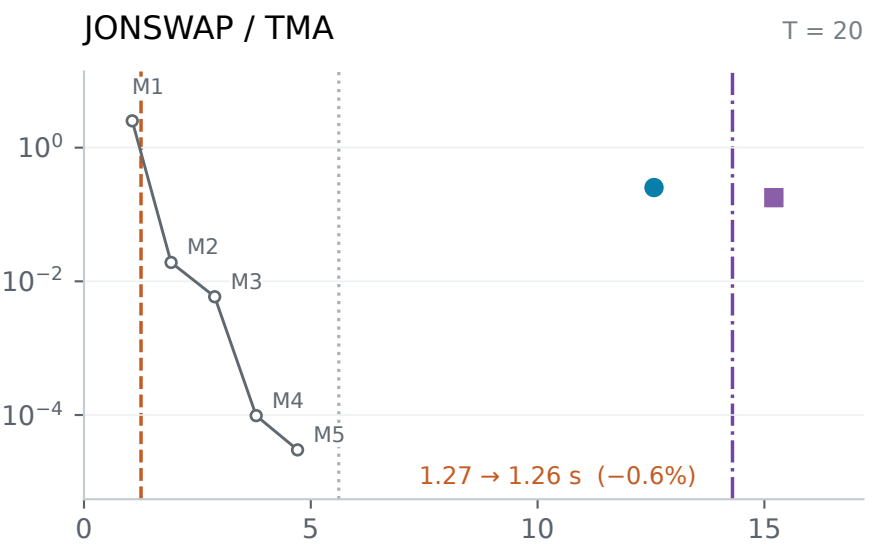

Median final surface error (%)

One rollout (seconds)

A100 · Modal — A100-SXM4-80GB

Classical
Neural — small (saved accuracy)
Neural — full (saved accuracy)
M6 — interleaved FFTs
New candidate optimized network
Spectral DNO — weight EMA (PyTorch)

Stokes — T = 20

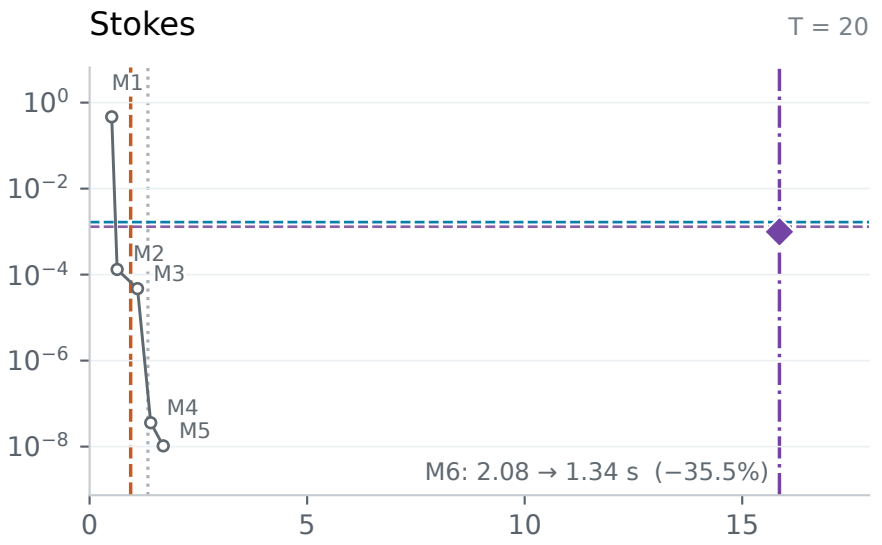

Tanaka — T = 200

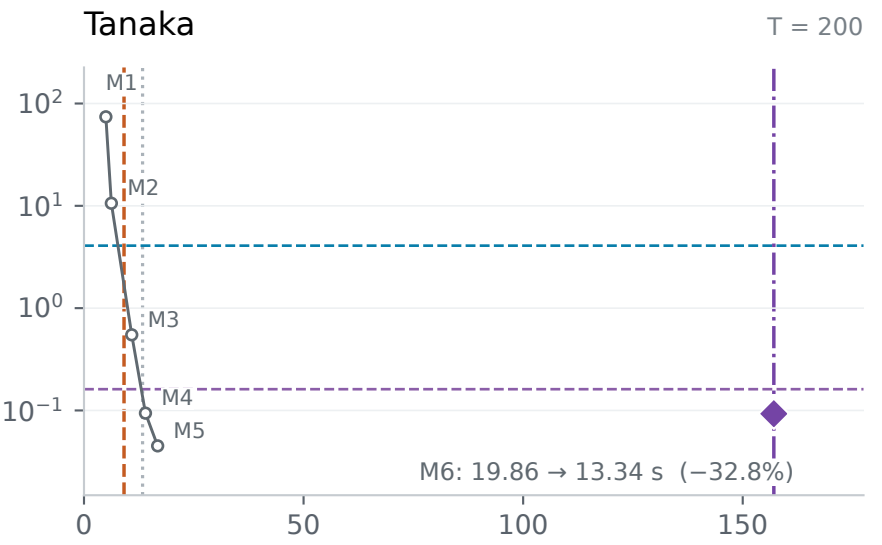

Benjamin–Feir — T = 200

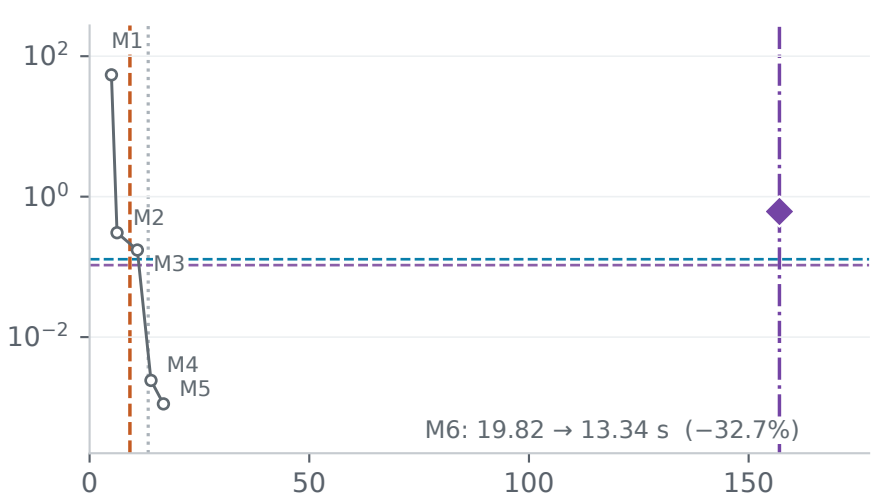

JONSWAP / TMA — T = 20

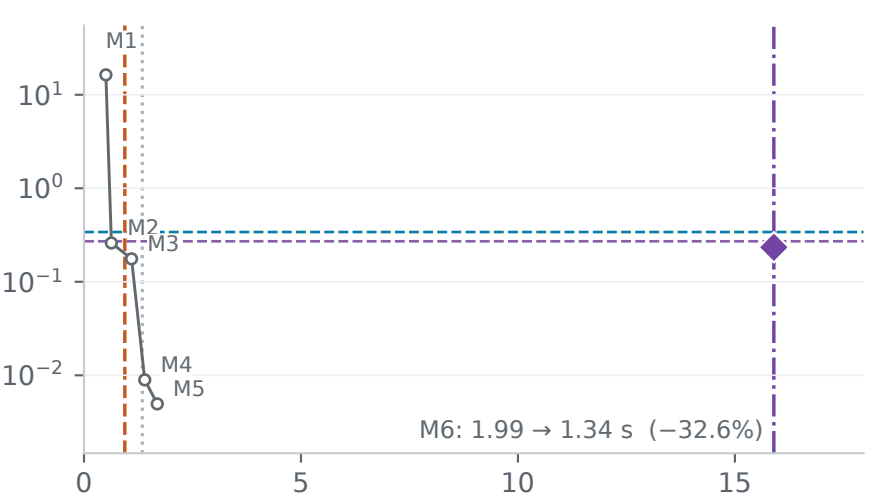

Final surface error (%)

One rollout (seconds)