


Single-rollout runtime vs error

—○— Classical ● Neural network — small model ■ Neural network — larger model M6 time

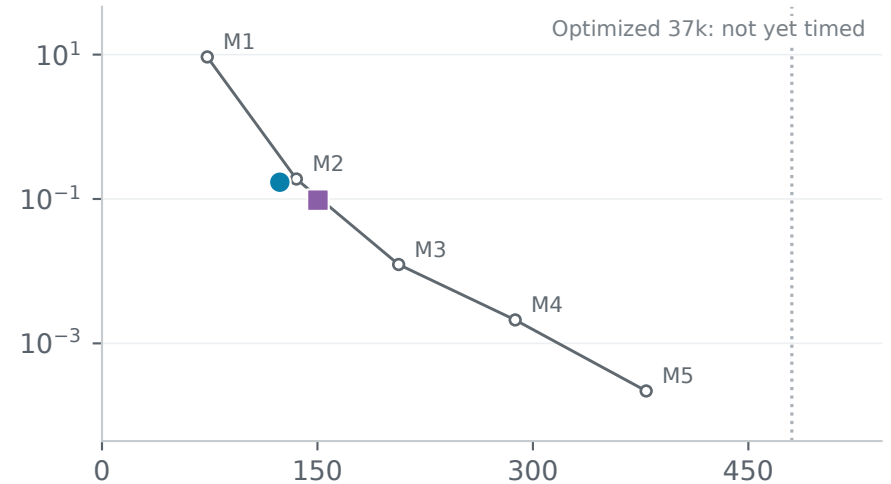
Stokes

T = 20



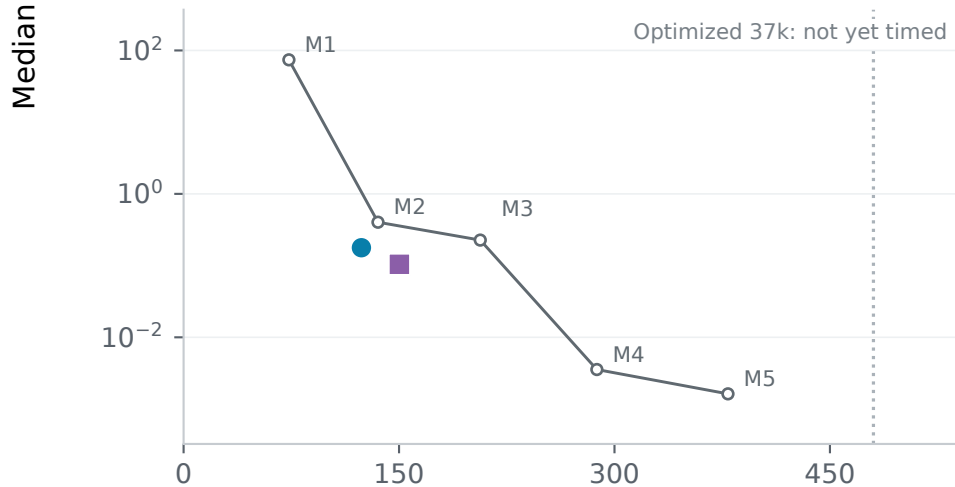
Tanaka

T = 200



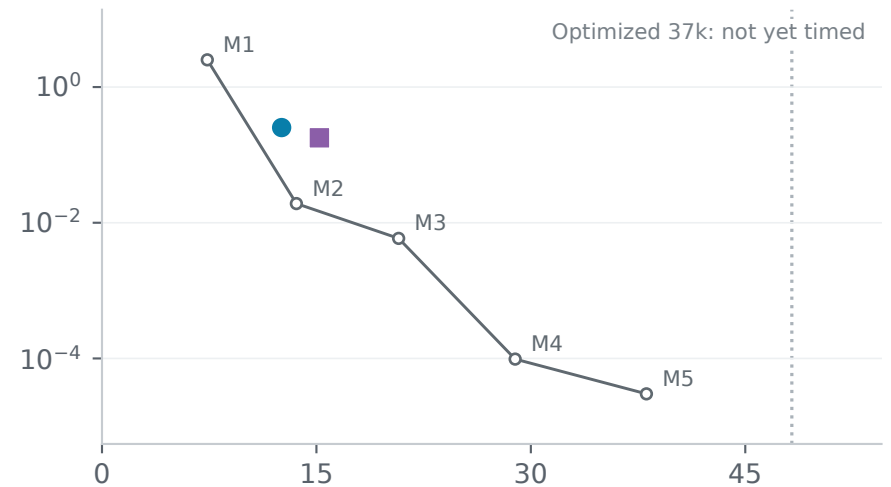
Benjamin-Feir

T = 200



JONSWAP / TMA

T = 20



One rollout (seconds)

Orange dashed line: measured optimized 37k runtime; median of 3 warmed single-rollout runs.
Existing points retain their 32-case error statistics. No 37k median accuracy or other-regime speed is assumed.