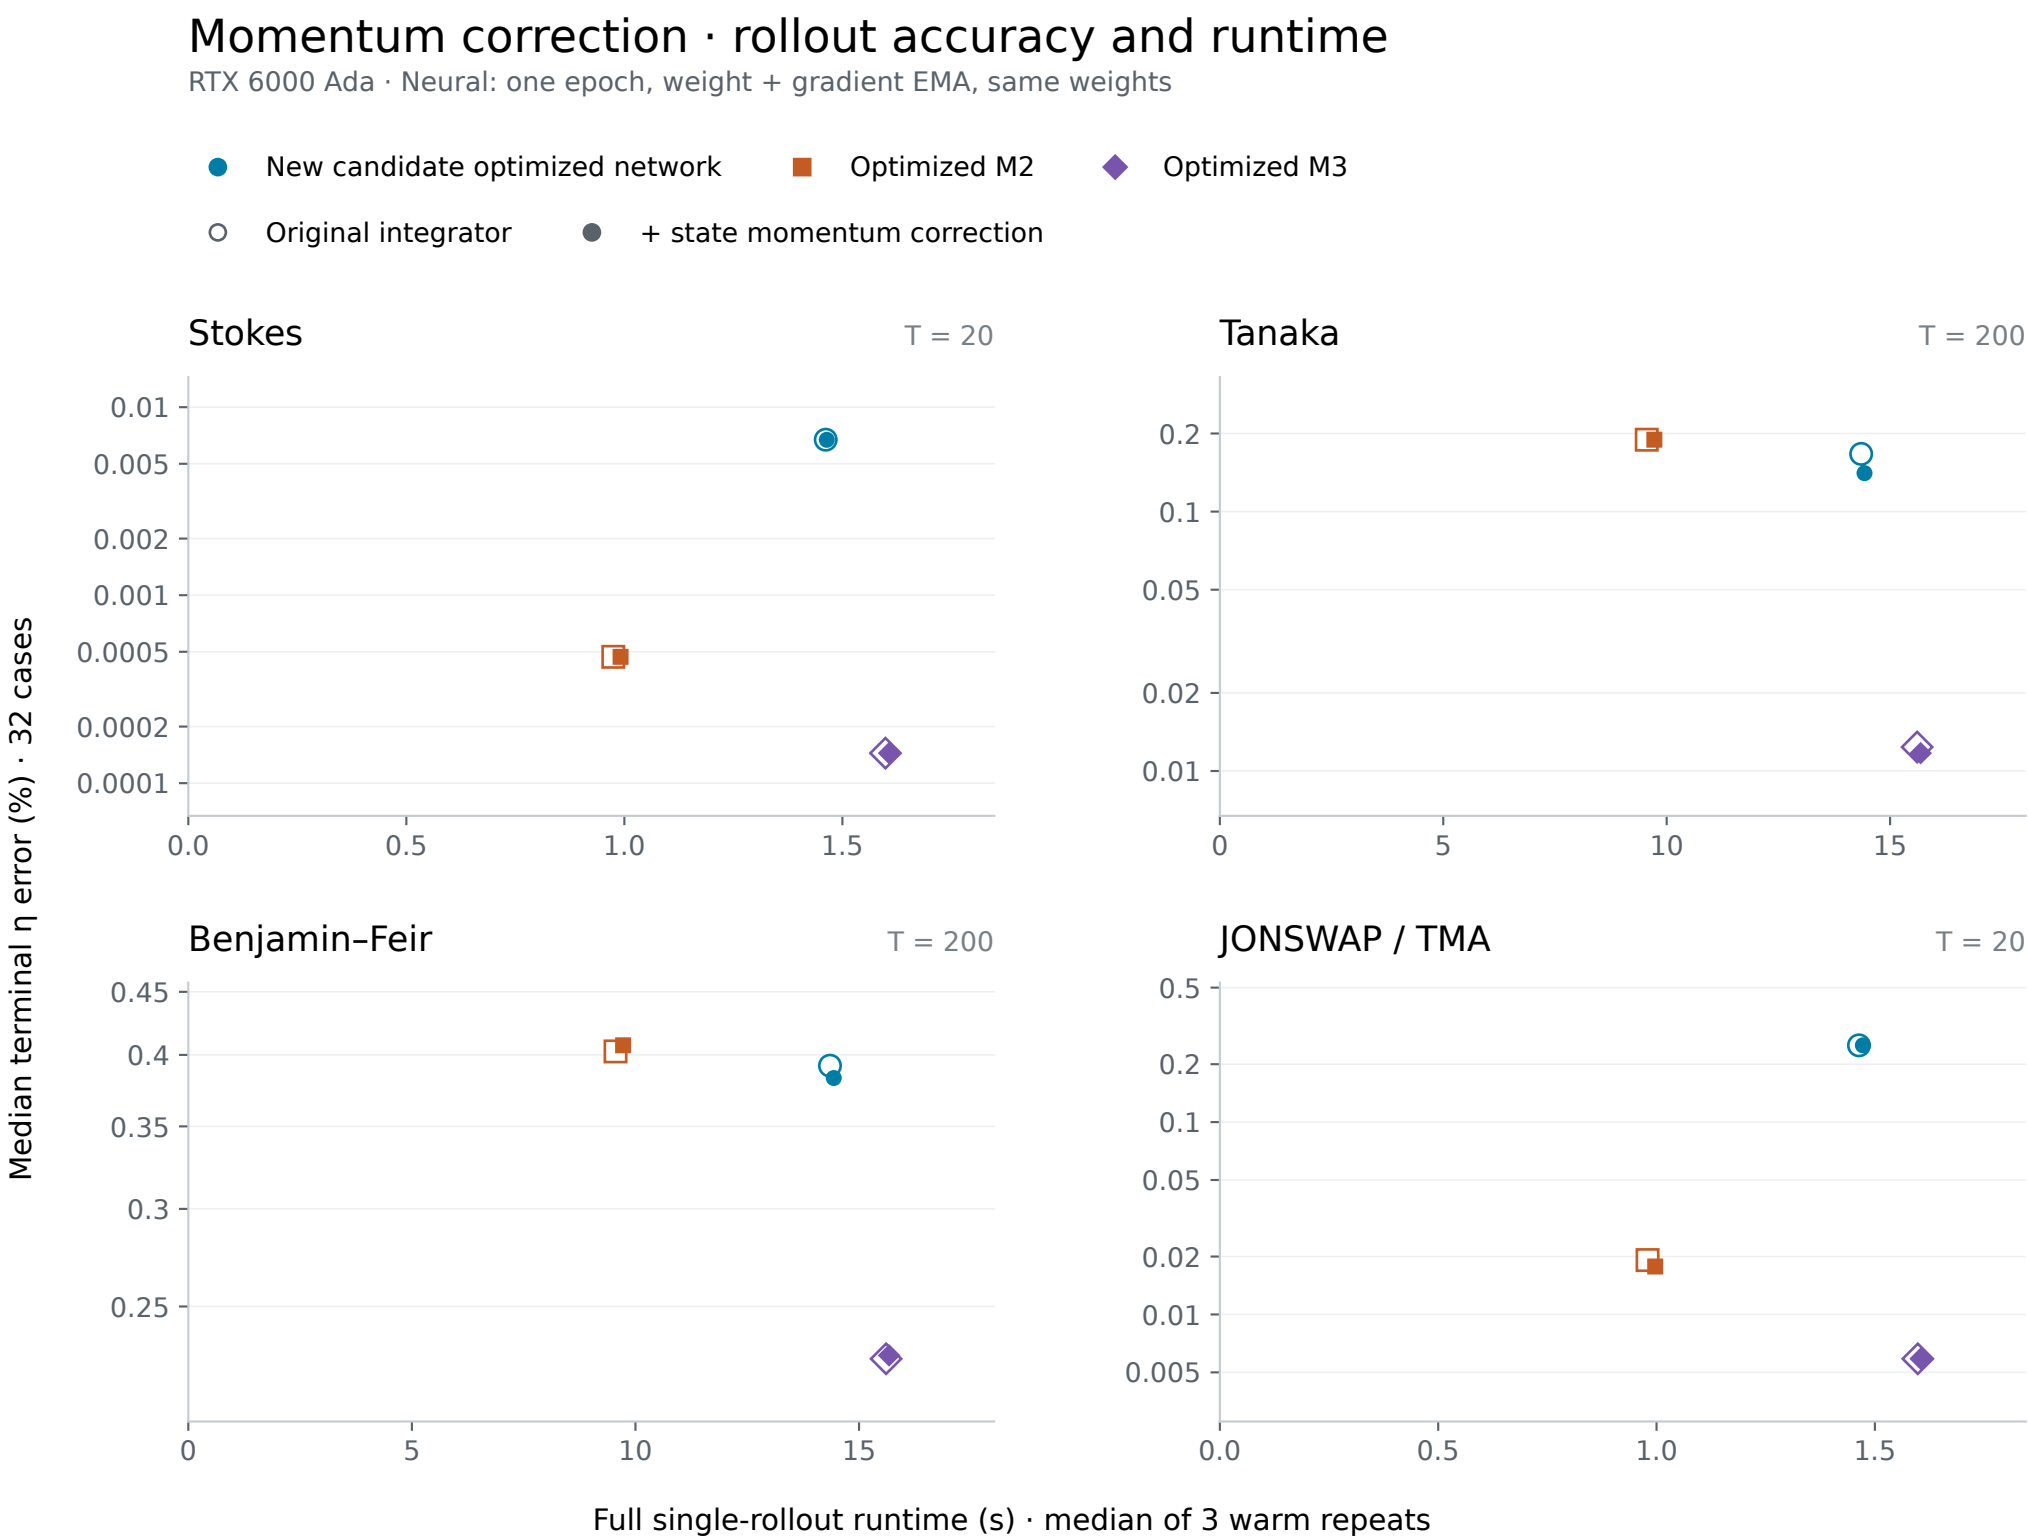
Momentum correction · rollout accuracy and runtime
RTX 6000 Ada · Neural: one epoch, weight + gradient EMA, same weights
New candidate optimized network
Optimized M2
Optimized M3
Original integrator
+ state momentum correction
Stokes
T = 20
Tanaka
T = 200
Benjamin–Feir
T = 200
JONSWAP / TMA
T = 20
Median terminal η error (%) · 32 cases
Full single-rollout runtime (s) · median of 3 warm repeats