Figure 12d · Optimized hero inference

RTX 6000 Ada · 32-case median error · Single-rollout timing · Same epoch-37 weight EMA
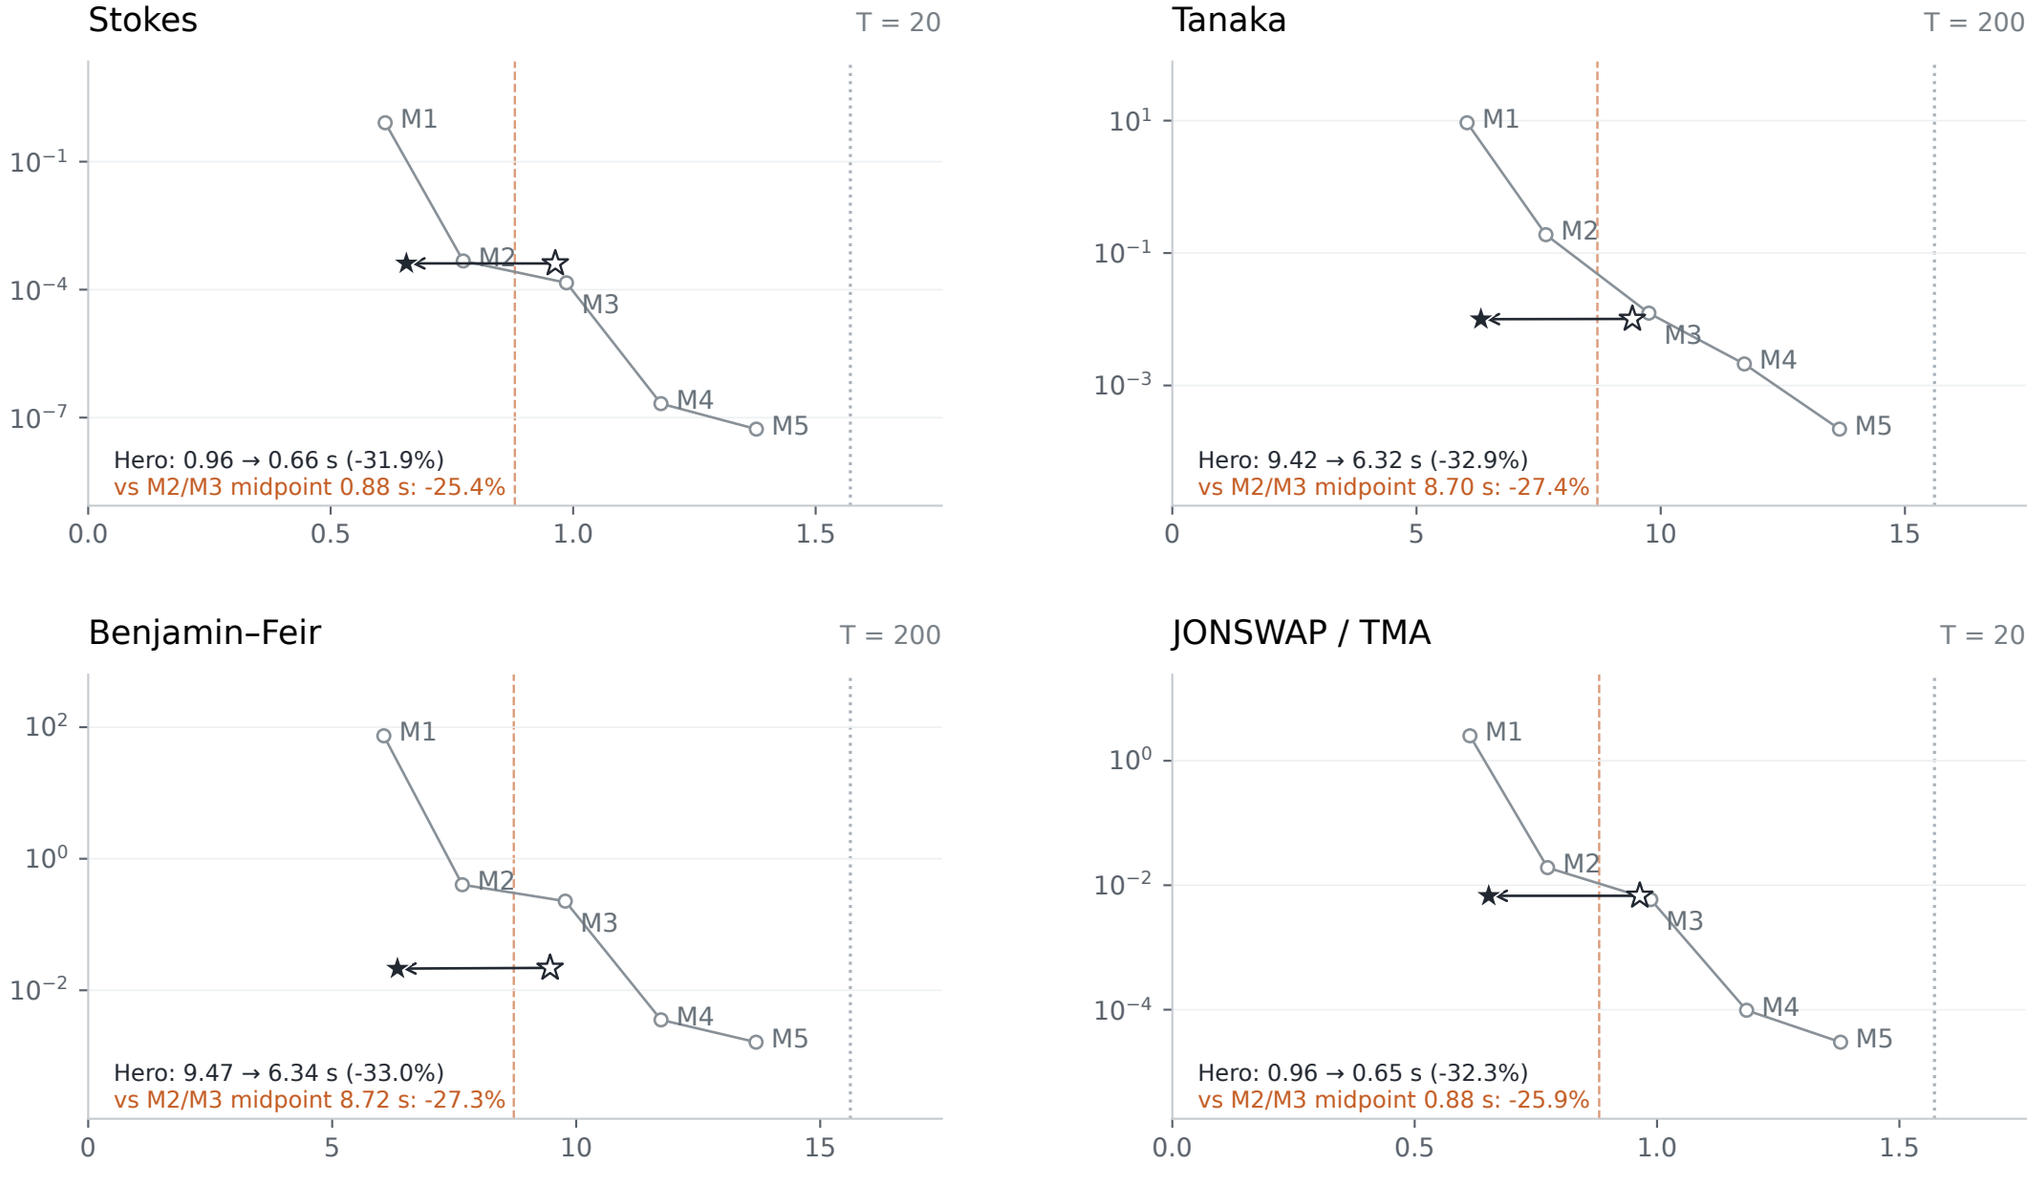
Classical
M6 time
M2/M3 runtime midpoint
Trained hero · epoch 37 · published inference
Trained hero · epoch 37 · optimized inference